

Neural Networks

Introduction, multilayer perceptron, optimization techniques

Machine Learning and Data Mining, 2021

Artem Maevskiy

National Research University Higher School of Economics



LAMBD A • HSE



November 10, 2021

From linear model to a neural network





Linear models + feature expansion recap

- ▶ Recall how, for linear models, we introduced **new features** to make the model **more powerful**
- ▶ Finding good features (aka feature engineering) is a **highly non-trivial task**




Linear models + feature expansion recap

- ▶ Recall how, for linear models, we introduced **new features** to make the model **more powerful**
- ▶ Finding good features (aka feature engineering) is a **highly non-trivial task**




Can we automate feature engineering? ☺

Idea: add another model




- Add another model

Idea: add another model



- ▶ Add another model
- ▶ Train everything simultaneously
 - Can use gradient descent if both models are **differentiable**


Idea: add another model




- ▶ Add another model
- ▶ Train everything simultaneously
 - Can use gradient descent if both models are **differentiable**

Note: stacking models like this likely makes the problem non-convex
⇒ **no convergence guarantees**

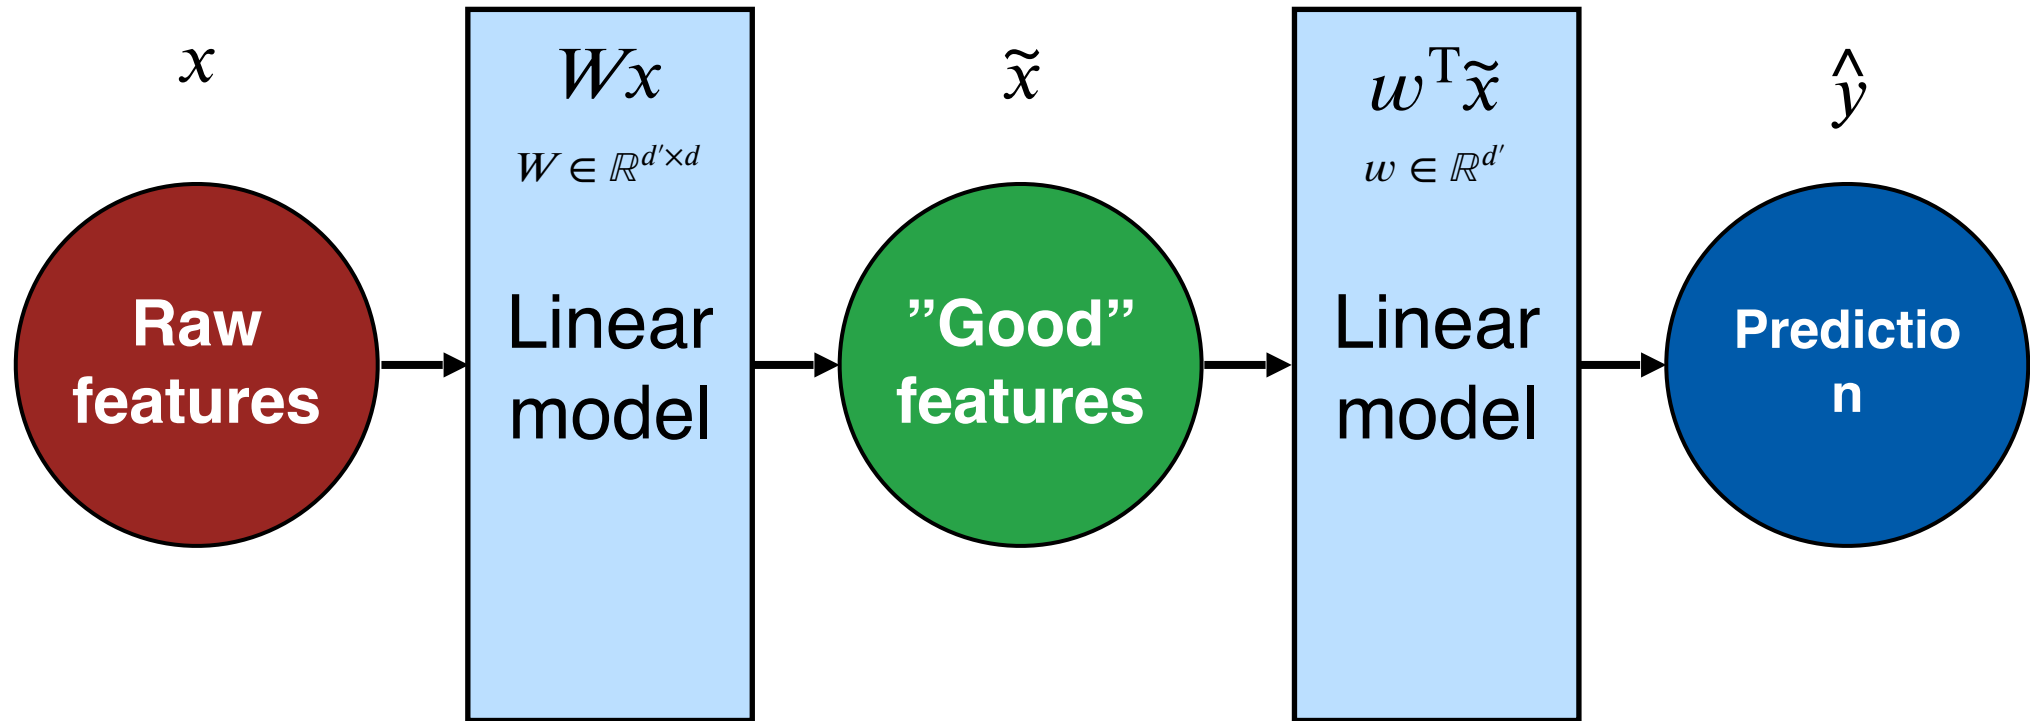
Can it be another linear model?



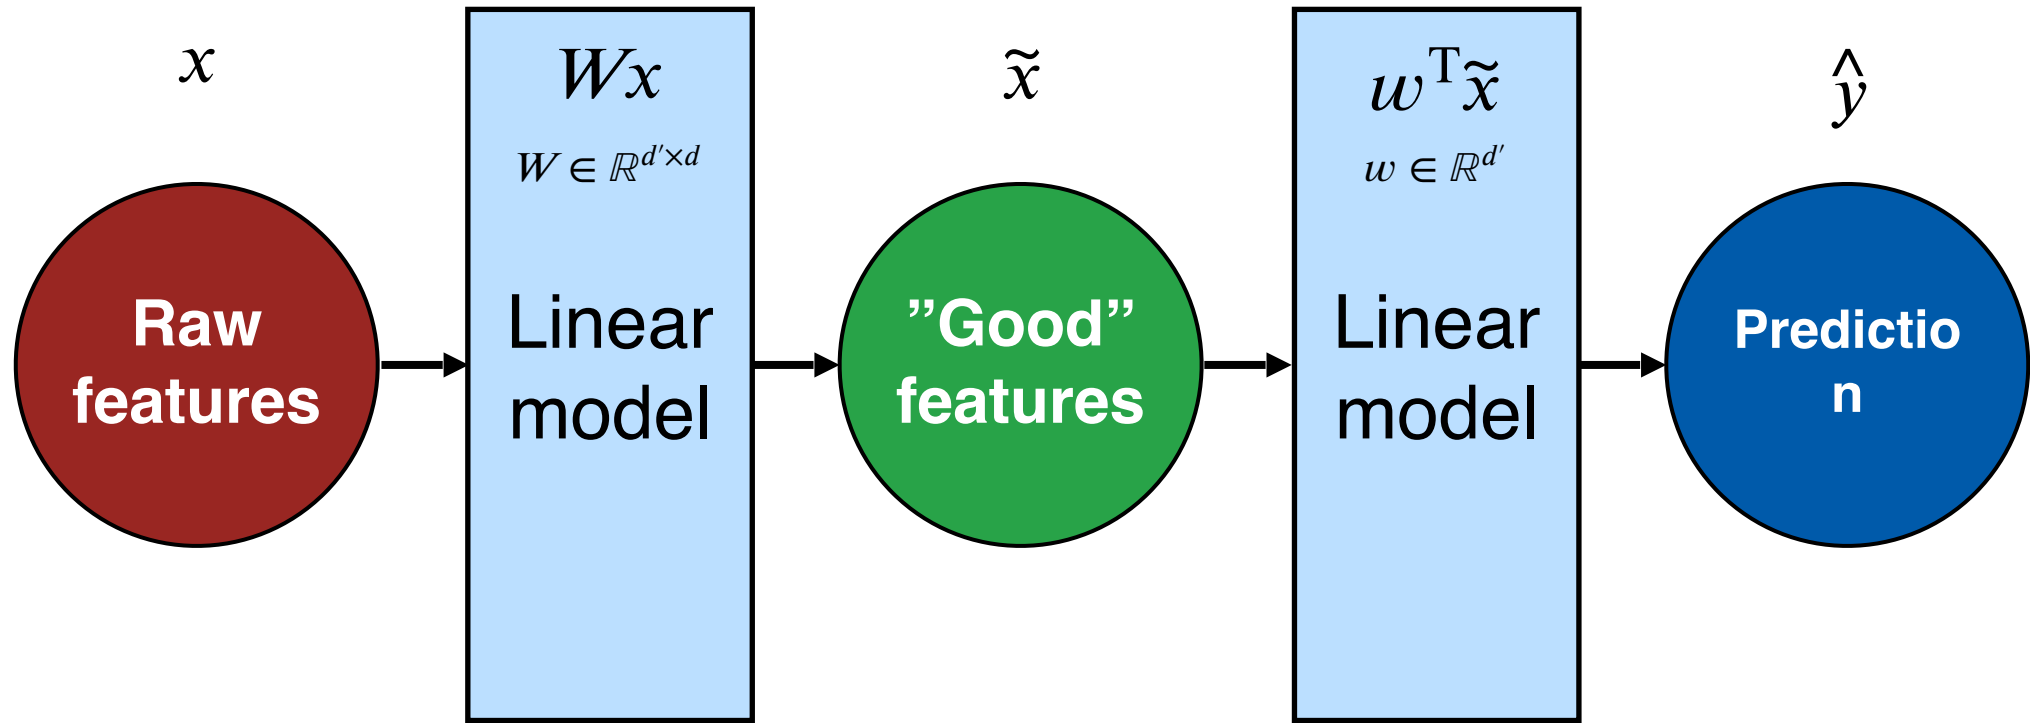
Can it be another linear model?



Can it be another linear model?



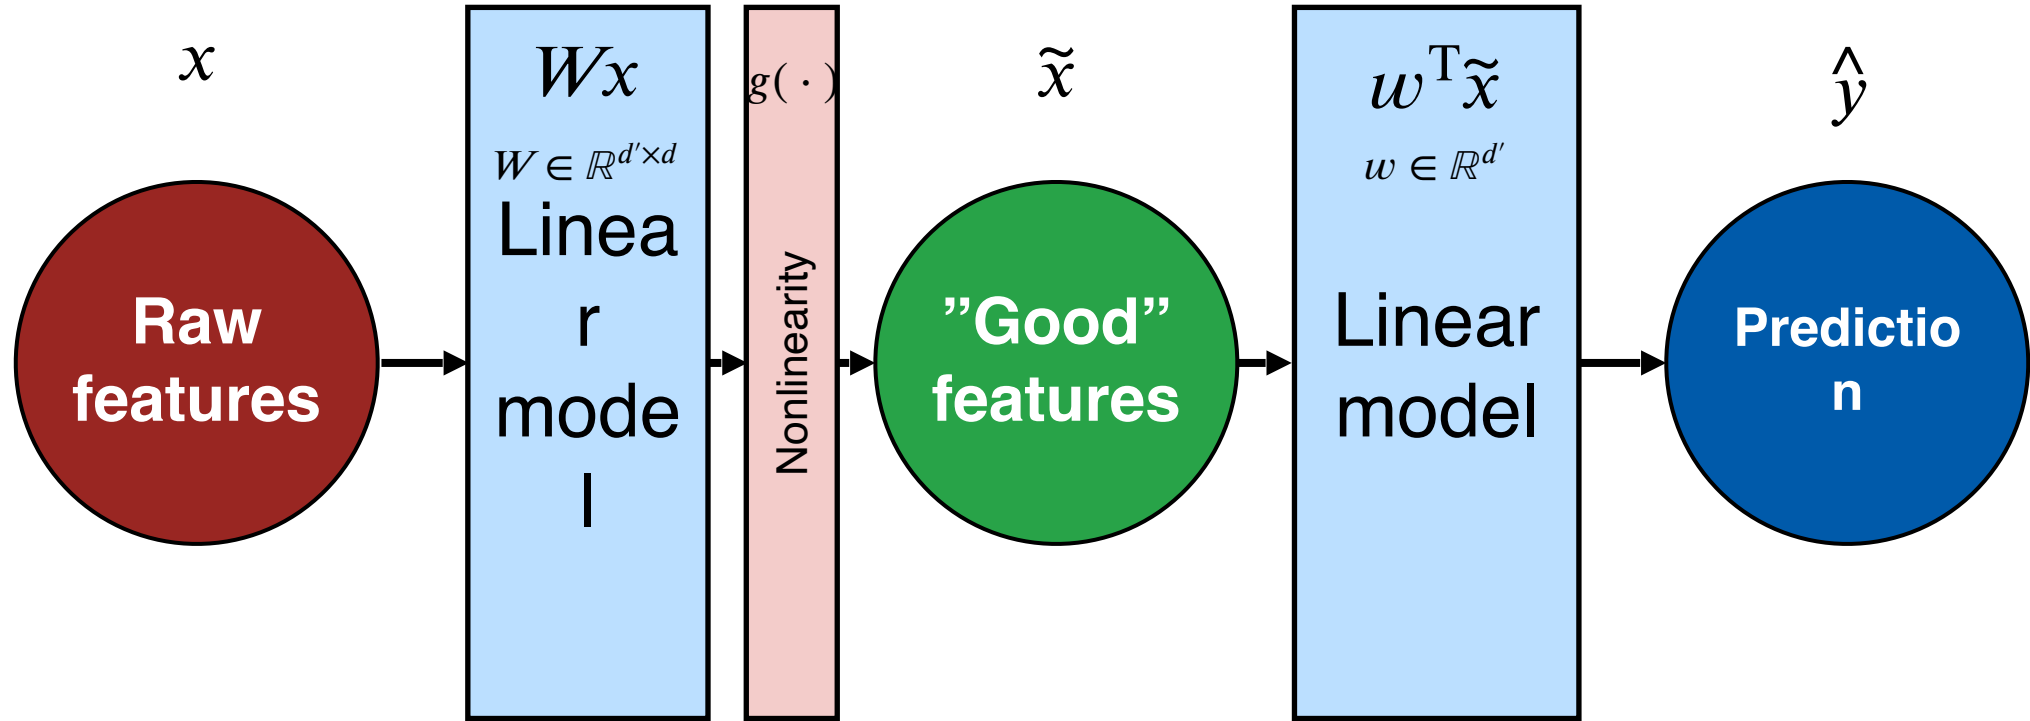
Can it be another linear model?



$$\hat{y} = w^T \tilde{x} = w^T (Wx) = (w^T W) x = w'^T x$$

– turns everything into just a single linear model

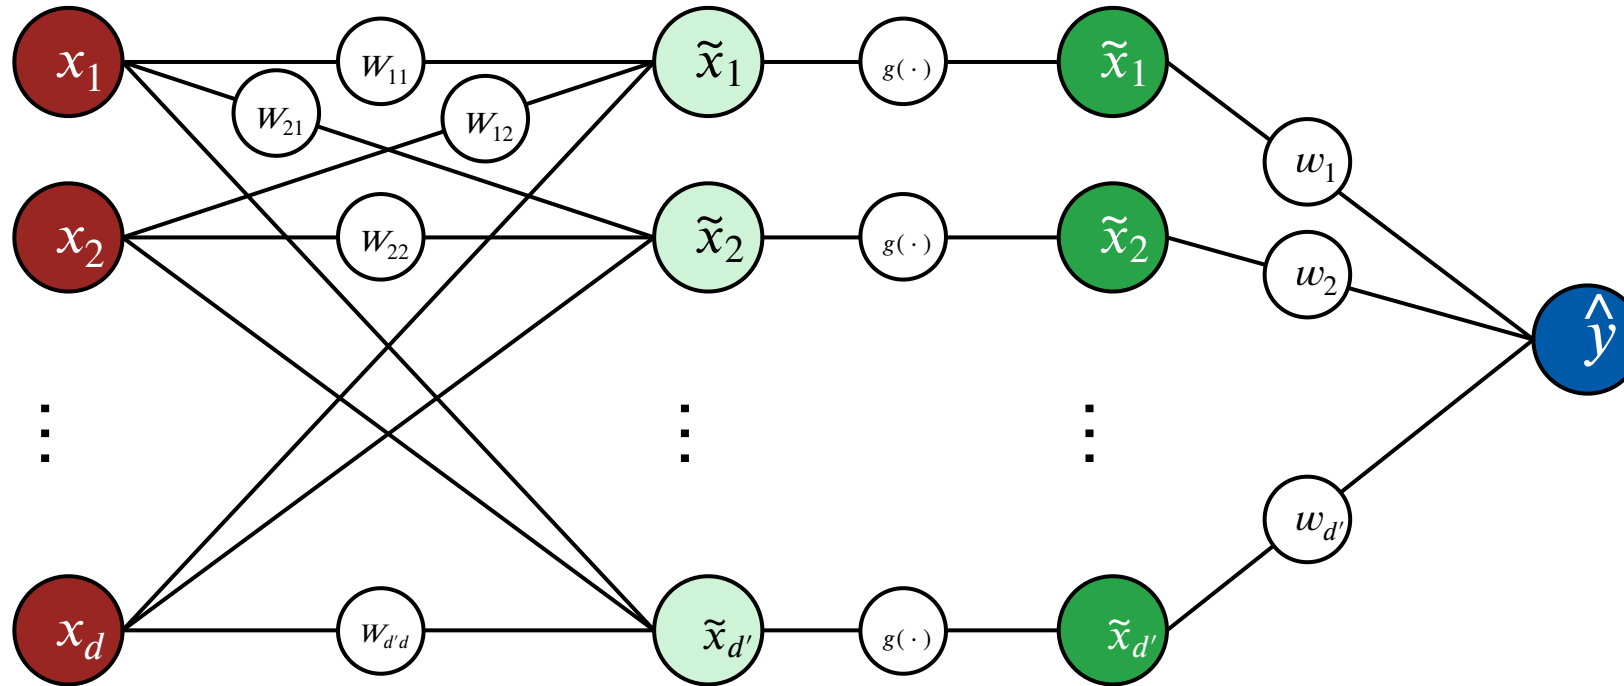
Fix: just introduce a nonlinearity



$$\hat{y} = w^T \tilde{x} = w^T g(Wx)$$

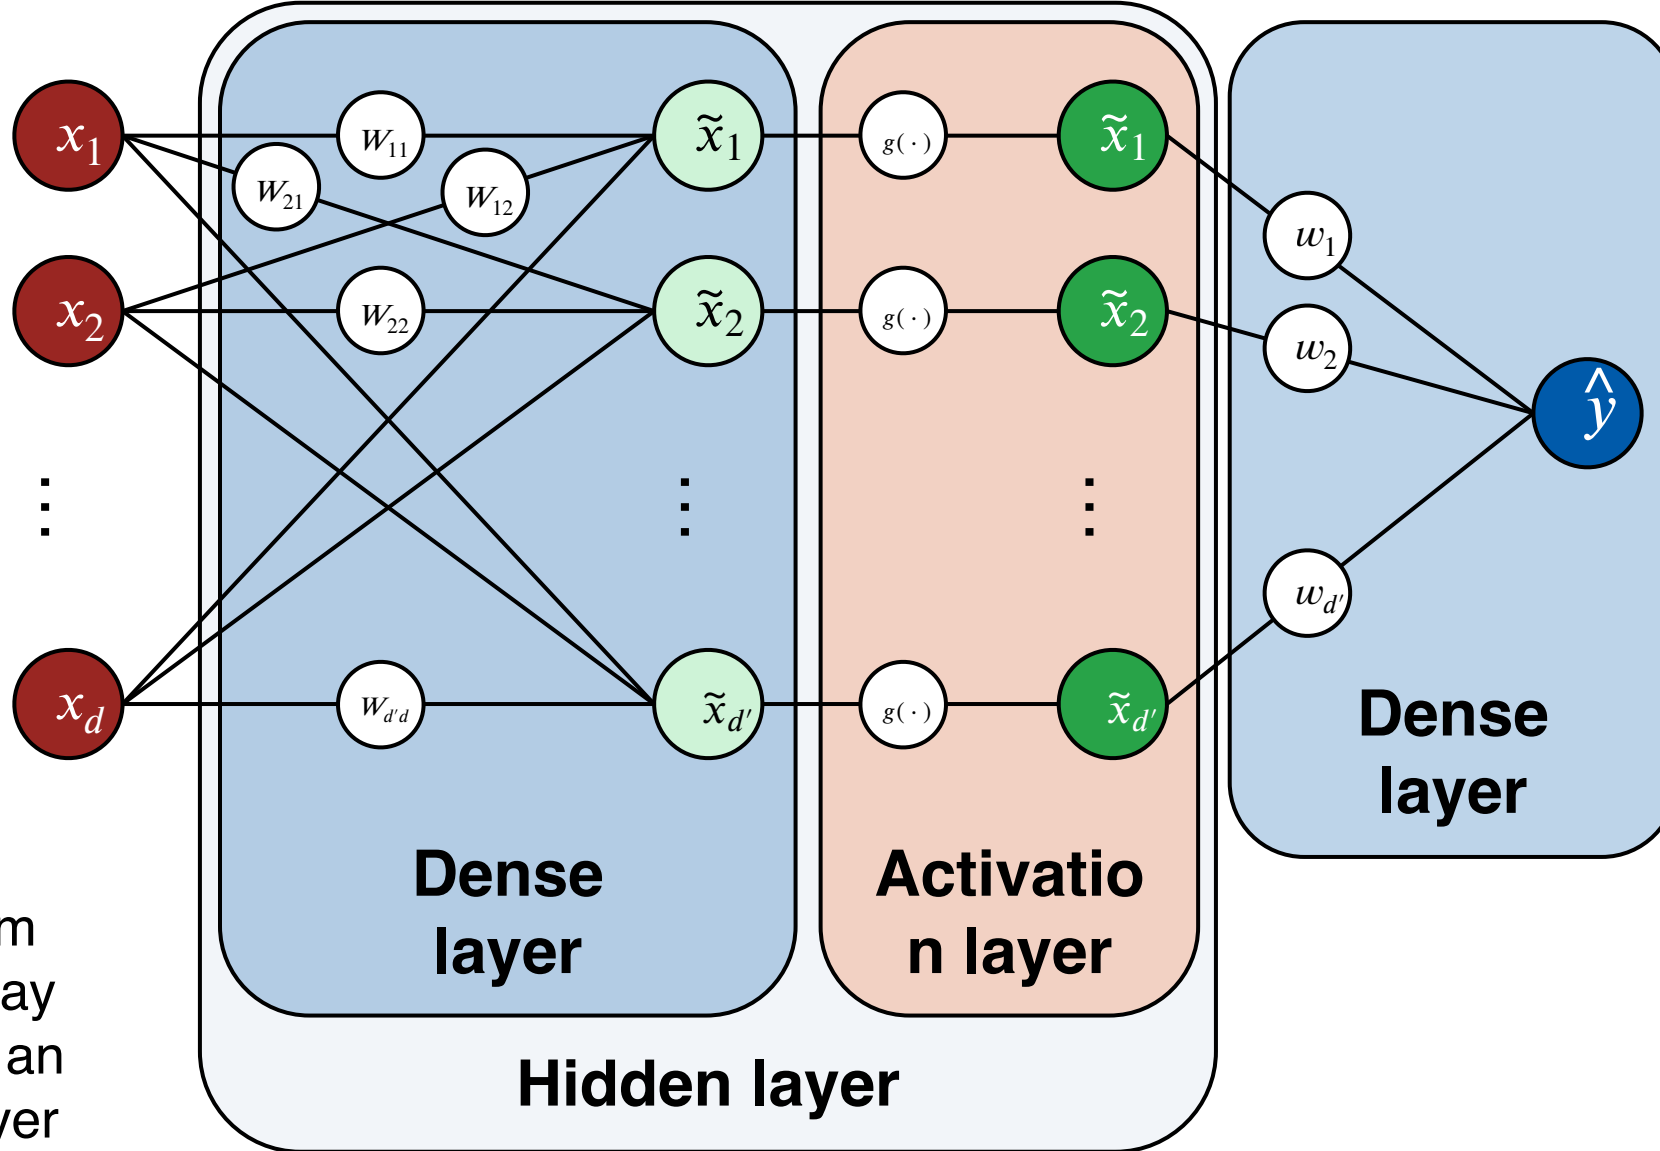
$g(\cdot)$ – some **nonlinear** scalar function (applied elementwise)

In greater detail



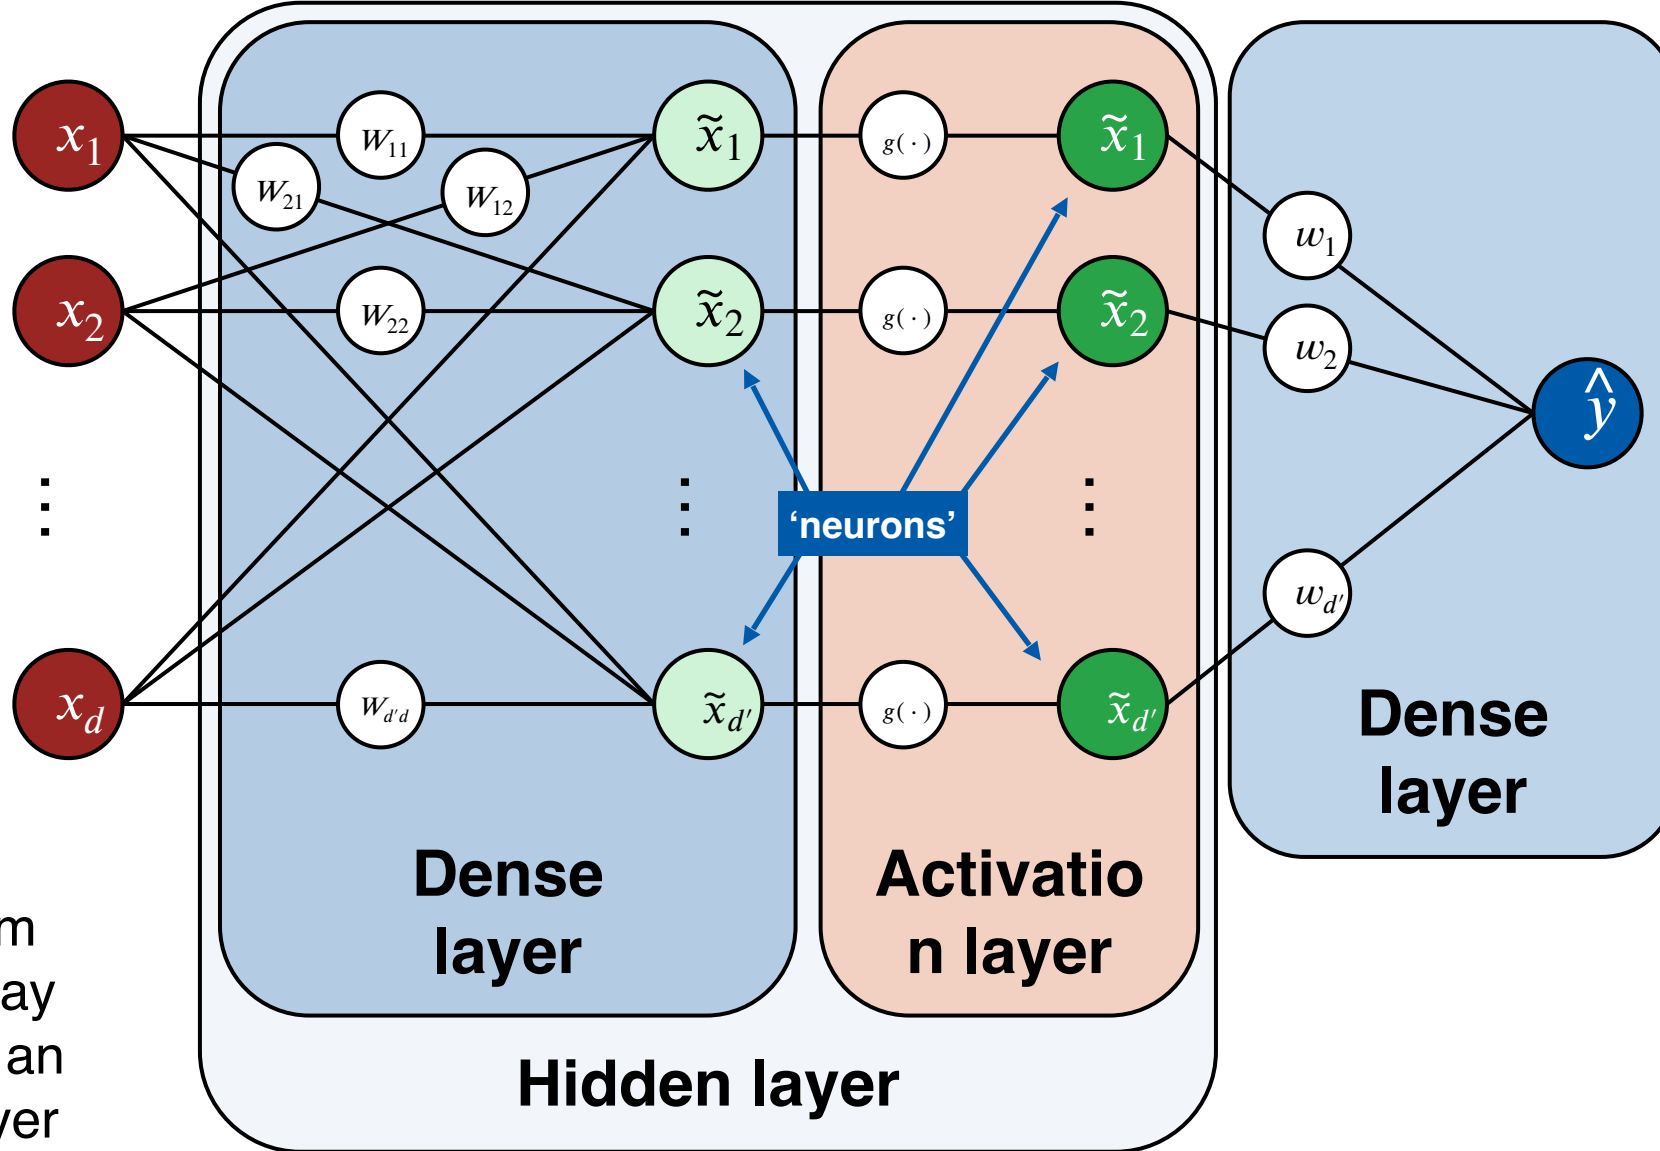
$$\hat{y} = w^T \tilde{x} = w^T g(Wx) = \sum_j \left[w_j g \left(\sum_i W_{ji} x_i \right) \right]$$

Some terminology



Note: the term “activation” may also stand for an output of a layer

Some terminology




Note: the term “activation” may also stand for an output of a layer

Activation functions




Activation functions



$$\text{ReLU}(x) = \max(0, x)$$


Activation functions



$$\text{ReLU}(x) = \max(0, x)$$

$$\text{ELU}(x) = \begin{cases} x & x \geq 0 \\ e^x - 1 & x < 0 \end{cases}$$

Activation functions

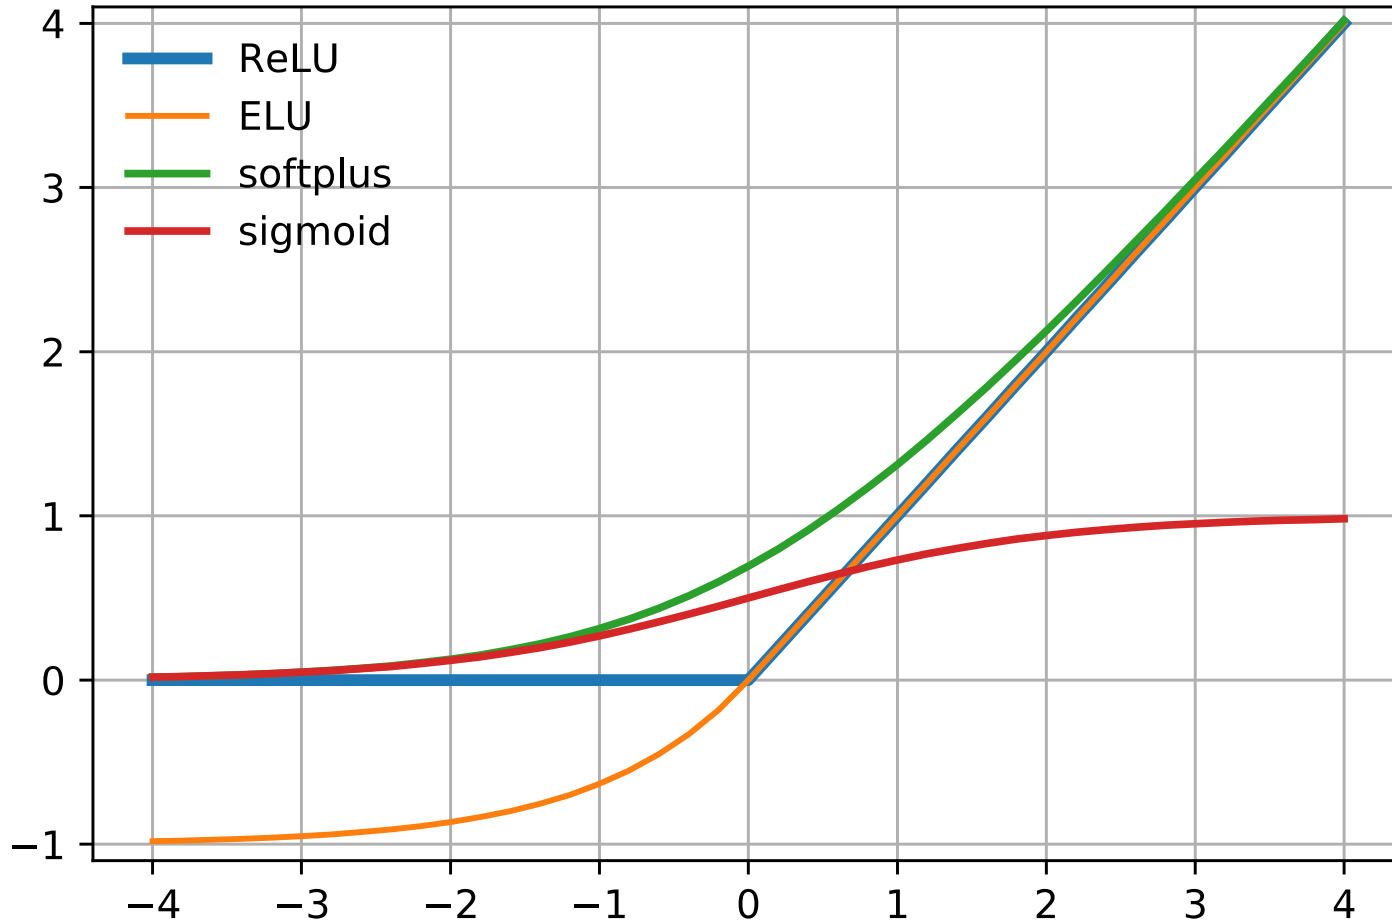


$$\text{ReLU}(x) = \max(0, x)$$

$$\text{ELU}(x) = \begin{cases} x & x \geq 0 \\ e^x - 1 & x < 0 \end{cases}$$

$$\text{softplus}(x) = \log(1 + e^x)$$

Activation functions




$$\text{ReLU}(x) = \max(0, x)$$

$$\text{ELU}(x) = \begin{cases} x & x \geq 0 \\ e^x - 1 & x < 0 \end{cases}$$

$$\text{softplus}(x) = \log(1 + e^x)$$

$$\text{sigmoid}(x) = \frac{1}{1 + e^{-x}}$$

Activation functions



$$\text{ReLU}(x) = \max(0, x)$$

$$\text{ELU}(x) = \begin{cases} x & x \geq 0 \\ e^x - 1 & x < 0 \end{cases}$$

$$\text{softplus}(x) = \log(1 + e^x)$$

$$\text{sigmoid}(x) = \frac{1}{1 + e^{-x}}$$


$$\text{tanh}(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

Universal approximator




Universal approximator

- ▶ Just a single hidden layer with a nonlinearity makes this model a universal approximator




Universal approximator

- ▶ Just a single hidden layer with a nonlinearity makes this model a universal approximator
 - any function can be approximated **arbitrarily close** given wide enough hidden layer (large enough d')

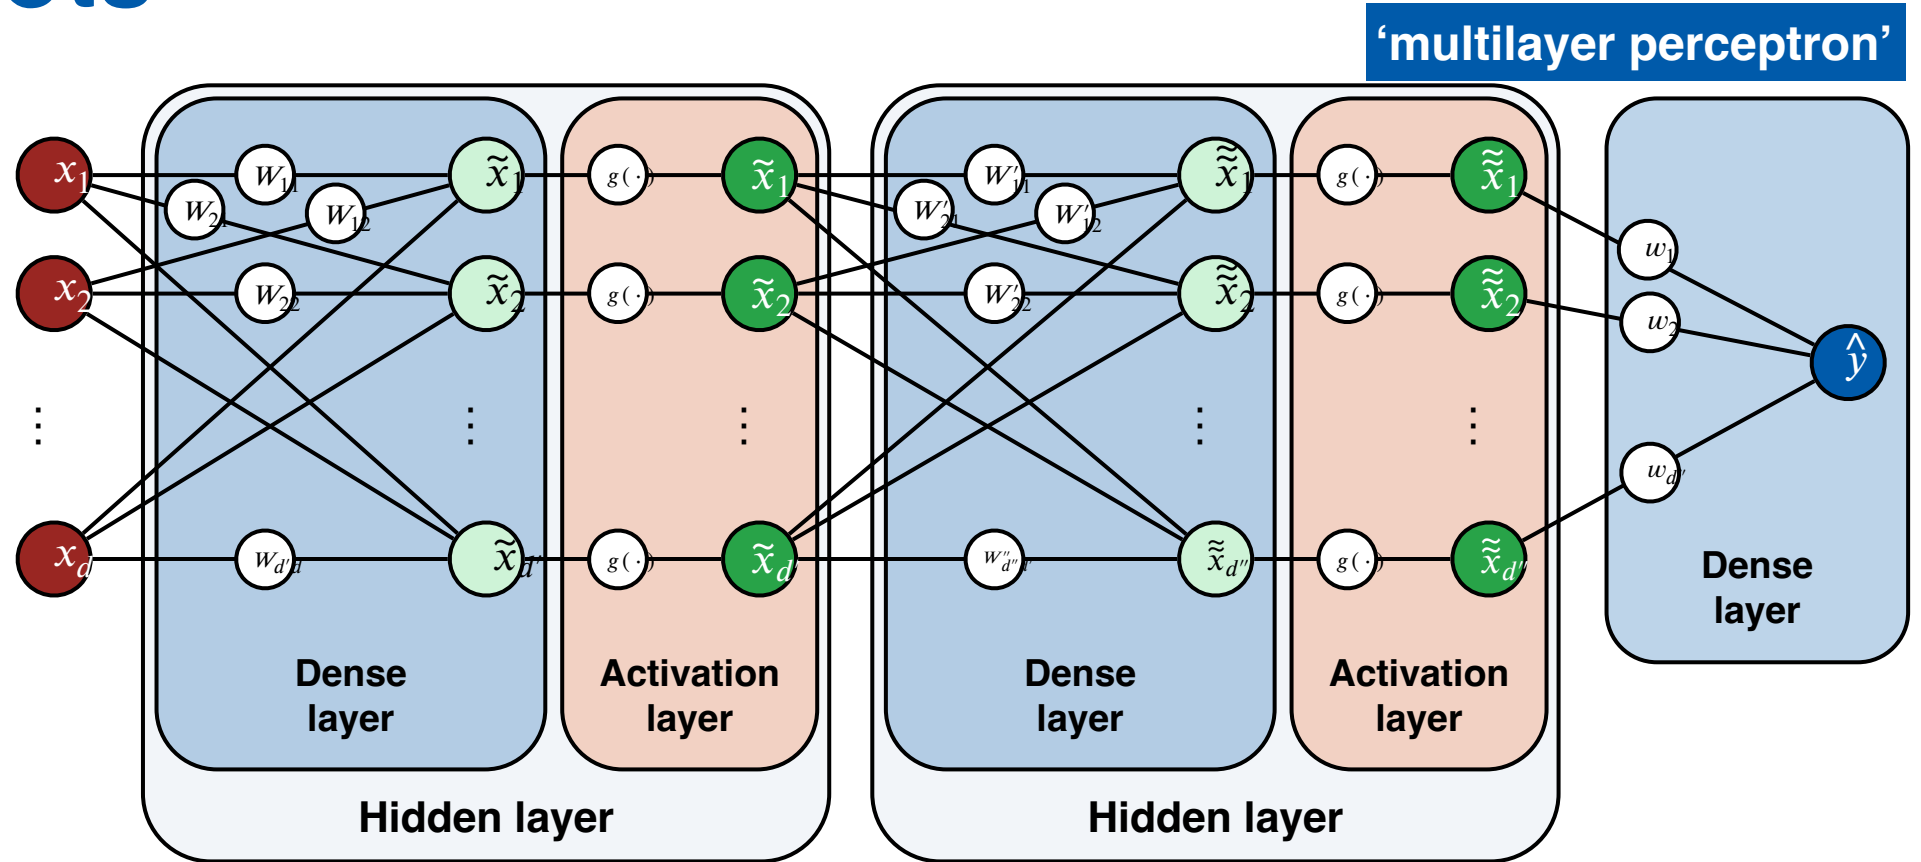


Universal approximator

- ▶ Just a single hidden layer with a nonlinearity makes this model a universal approximator
 - any function can be approximated **arbitrarily close** given wide enough hidden layer (large enough d')
 - Note: in practice we might not be able to find this approximation
 - e.g. due to heavily non-convex loss function, infeasibly large d' , overfitting



Deeper nets




- In practice, stacking more hidden layers often reduces the number of neurons required to represent a given function

Backpropagation



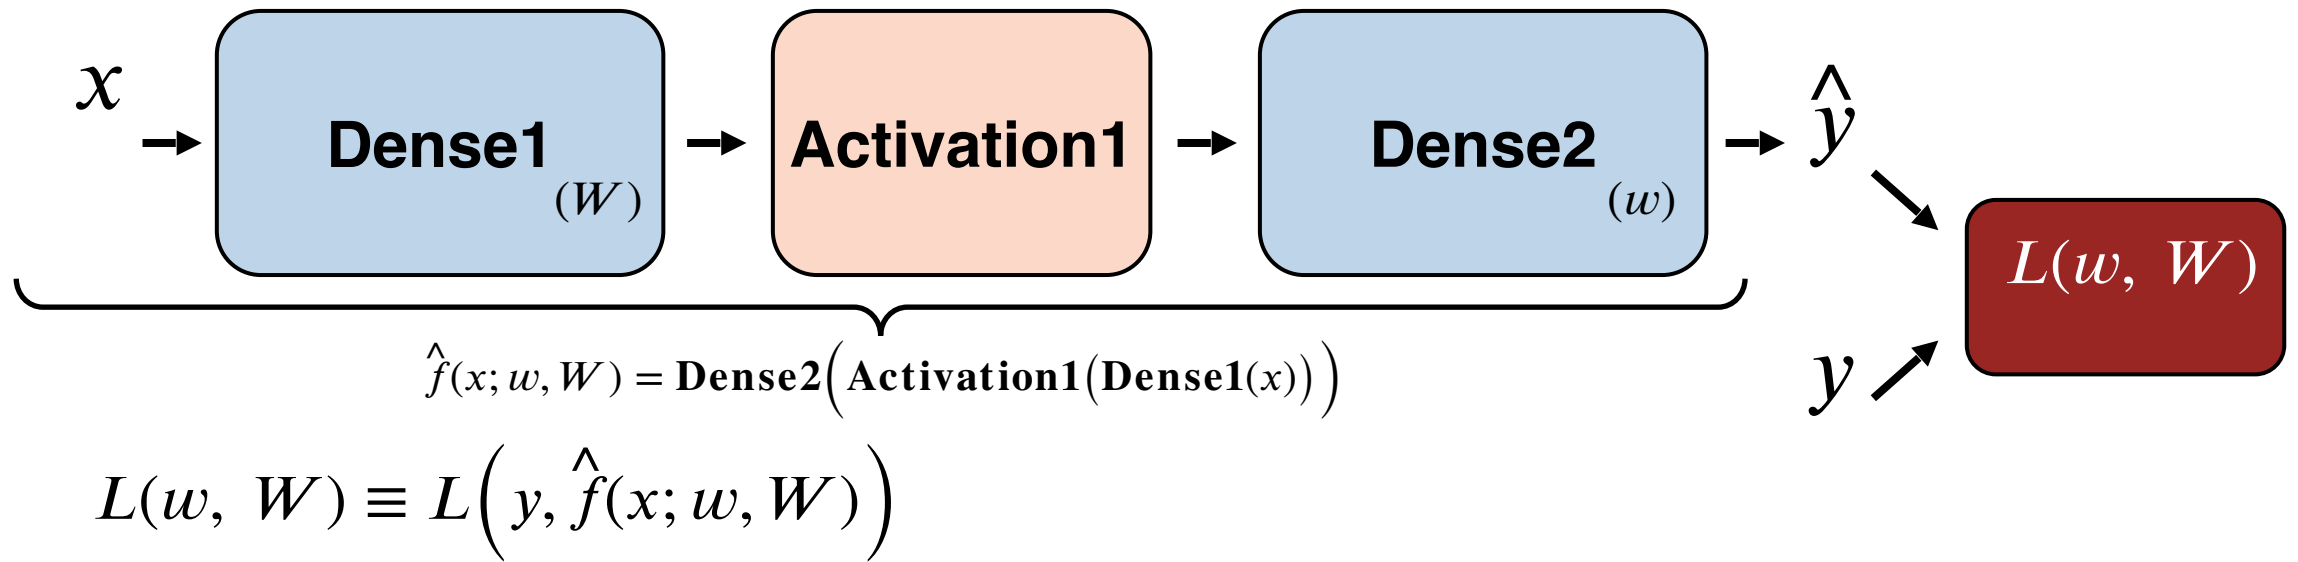
Loss function



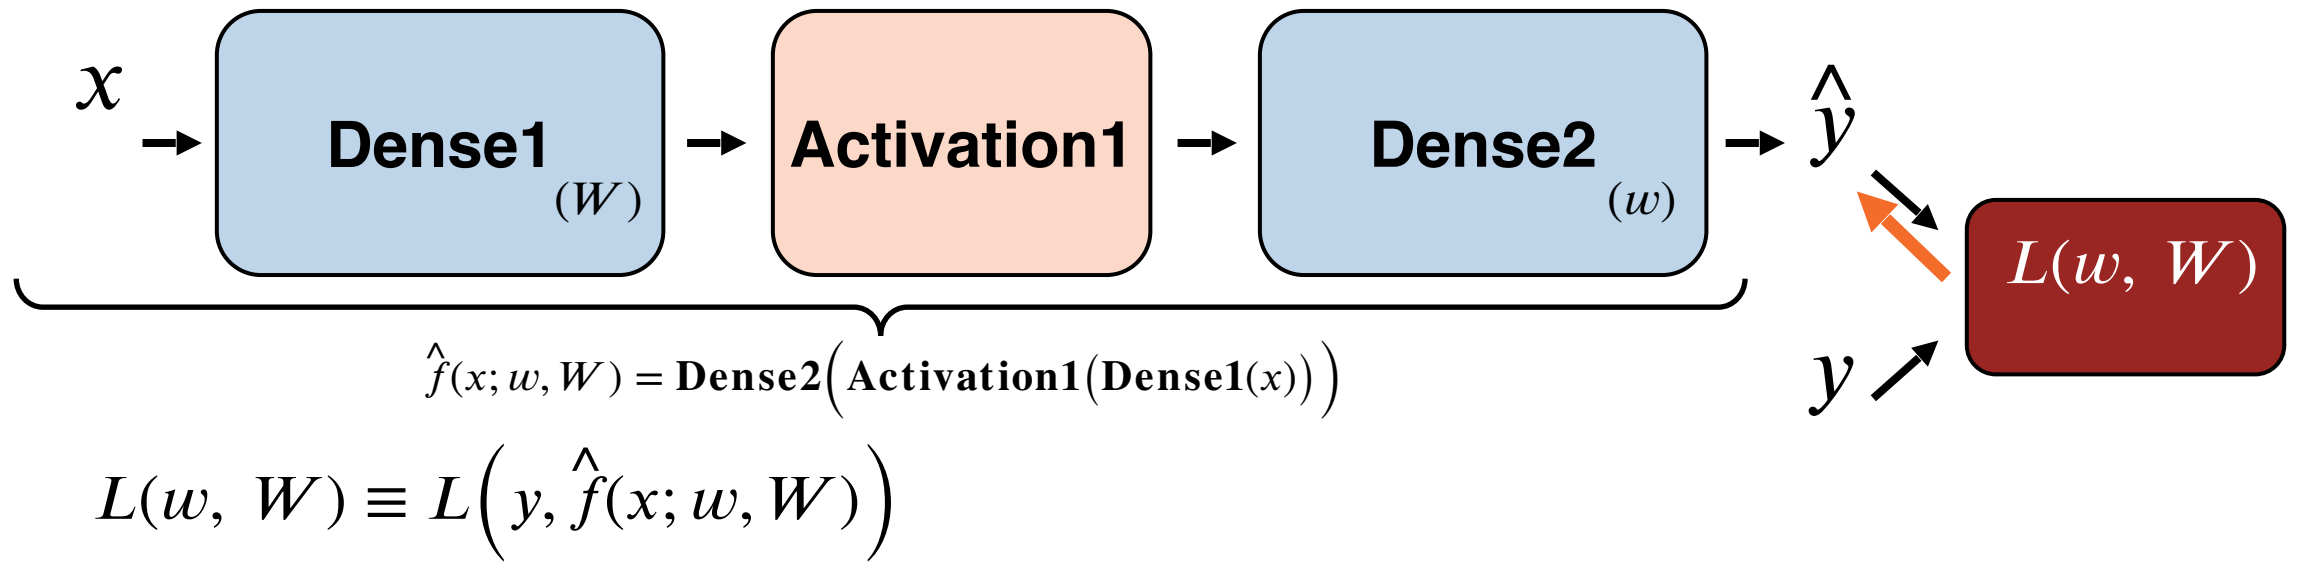
- E.g. mean squared error:

$$L = \frac{1}{N} \sum_{i=1 \dots N} \left(y_i - w^T g(Wx_i) \right)^2$$

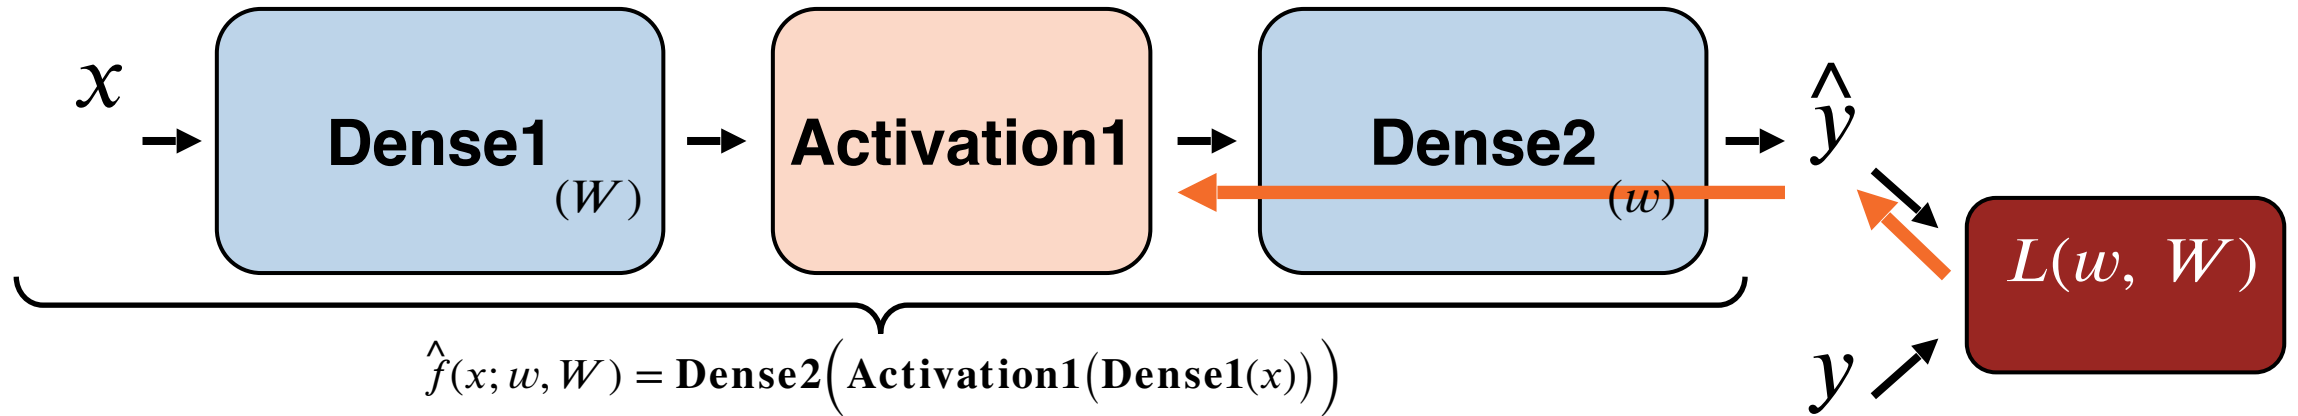
Calculating derivatives



Calculating derivatives

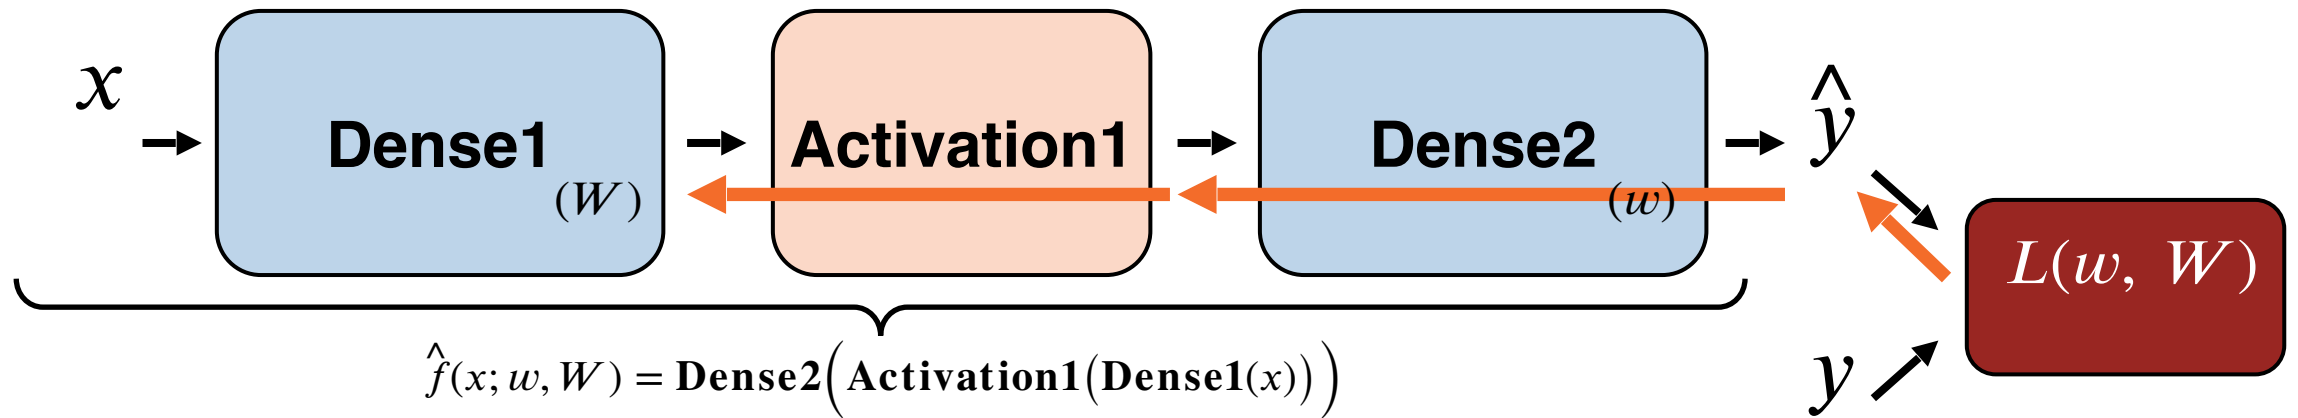


Calculating derivatives



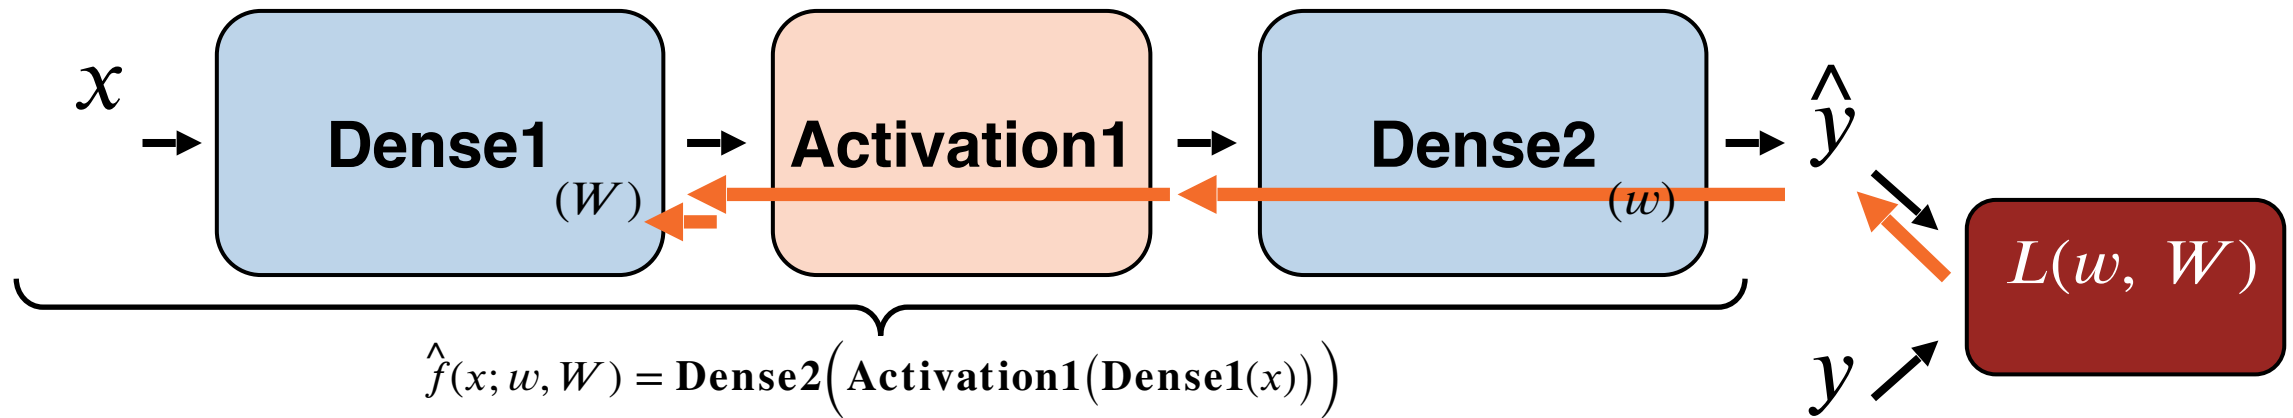
$$L(w, W) \equiv L\left(y, \hat{f}(x; w, W)\right)$$

Calculating derivatives



$$L(w, W) \equiv L\left(y, \hat{f}(x; w, W)\right)$$

Calculating derivatives

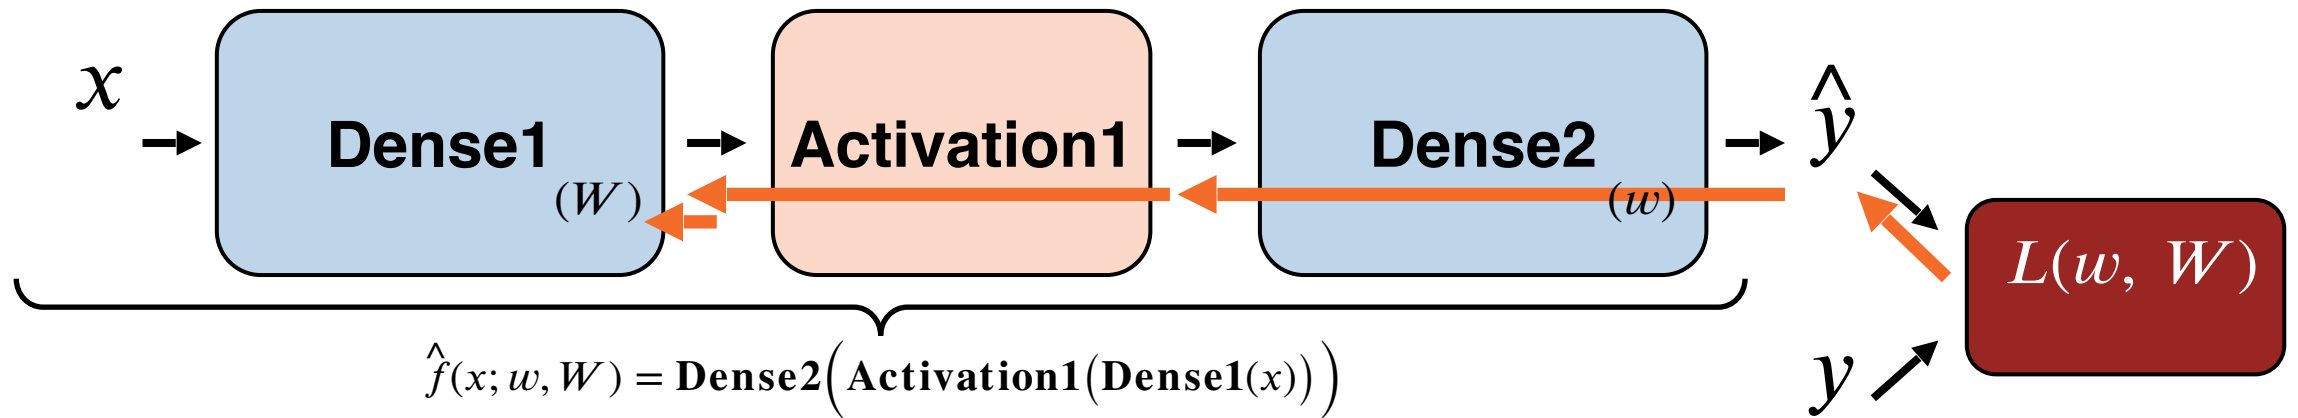


$$L(w, W) \equiv L\left(y, \hat{f}(x; w, W)\right)$$

$$\frac{\partial L}{\partial W} = \frac{\partial L(y, \hat{f})}{\partial \hat{f}} \cdot \frac{\partial \hat{f}}{\partial W} = \frac{\partial L(y, \hat{f})}{\partial \hat{f}} \cdot \frac{\partial \text{Dense2}}{\partial \text{Activation1}} \cdot \frac{\partial \text{Activation1}}{\partial \text{Dense1}} \cdot \frac{\partial \text{Dense1}}{\partial W}$$

- Backpropagation algorithm \approx applying the chain rule
 - The actual algorithm states how to do it efficiently

Calculating derivatives



$$L(w, W) \equiv L(y, \hat{f}(x; w, W))$$

$$\frac{\partial L}{\partial W} = \frac{\partial L(y, \hat{f})}{\partial \hat{f}} \cdot \frac{\partial \hat{f}}{\partial W} = \underbrace{\frac{\partial L(y, \hat{f})}{\partial \hat{f}}}_{\text{scalar}} \cdot \underbrace{\frac{\partial \text{Dense2}}{\partial \text{Activation1}}}_{d'\text{-vector}} \cdot \underbrace{\frac{\partial \text{Activation1}}{\partial \text{Dense1}}}_{d' \times d'\text{-matrix (diagonal)}} \cdot \underbrace{\frac{\partial \text{Dense1}}{\partial W}}_{d' \times d' \times d\text{-tensor} = d' \times d\text{-matrix}}$$


- Backpropagation algorithm \approx applying the chain rule
 - The actual algorithm states how to do it efficiently

Optimization techniques



How to optimize such functions?

Loss function landscape
(in 2 random directions around an optimum)




- ▶ No convergence guarantees for the stochastic gradient descent

<https://papers.nips.cc/paper/7875-visualizing-the-loss-landscape-of-neural-nets>

How to optimize such functions?

Loss function landscape
(in 2 random directions around an optimum)



- ▶ No convergence guarantees for the stochastic gradient descent
- ▶ There's a number of modifications to improve training


<https://papers.nips.cc/paper/7875-visualizing-the-loss-landscape-of-neural-nets>

SGD on mini-batches

► SGD:

- At each step k pick $l_k \in \{1, \dots, N\}$ at random, then update:

- $$\theta^{(k)} \leftarrow \theta^{(k-1)} - \eta \nabla_{\theta} \mathcal{L} \left(y_{l_k}, \hat{f}_{\theta}(x_{l_k}) \right) \mid \theta = \theta^{(k-1)}$$




SGD on mini-batches

► SGD:

- At each step k pick $l_k \in \{1, \dots, N\}$ at random, then update:
- $\theta^{(k)} \leftarrow \theta^{(k-1)} - \eta \nabla_{\theta} \mathcal{L} \left(y_{l_k}, \hat{f}_{\theta}(x_{l_k}) \right) \mid \theta = \theta^{(k-1)}$

► Mini-batch SGD:

- Shuffle the training set, then iterate through it in chunks (batches) of fixed size



SGD on mini-batches

► SGD:


- At each step k pick $l_k \in \{1, \dots, N\}$ at random, then update:

- $$\theta^{(k)} \leftarrow \theta^{(k-1)} - \eta \nabla_{\theta} \mathcal{L} \left(y_{l_k}, \hat{f}_{\theta}(x_{l_k}) \right) \mid \theta = \theta^{(k-1)}$$

► Mini-batch SGD:

- Shuffle the training set, then iterate through it in chunks (batches) of fixed size
- At each iteration evaluate the loss gradients on the given chunk B :

$$g = \sum_{i \in B} \nabla_{\theta} \mathcal{L} \left(y_i, \hat{f}_{\theta}(x_i) \right)$$




SGD on mini-batches

► SGD:


- At each step k pick $l_k \in \{1, \dots, N\}$ at random, then update:
- $\theta^{(k)} \leftarrow \theta^{(k-1)} - \eta \nabla_{\theta} \mathcal{L}\left(y_{l_k}, \hat{f}_{\theta}(x_{l_k})\right) \mid \theta = \theta^{(k-1)}$

► Mini-batch SGD:

- Shuffle the training set, then iterate through it in chunks (batches) of fixed size
- At each iteration evaluate the loss gradients on the given chunk B : $g = \sum_{i \in B} \nabla_{\theta} \mathcal{L}\left(y_i, \hat{f}_{\theta}(x_i)\right)$
- Update the model parameters: $\theta^{(k)} \leftarrow \theta^{(k-1)} - \eta \cdot g$




Momentum SGD



- Idea: introduce inertia (like a ball rolling down a hill)

$$m^{(k)} \longleftarrow \beta \cdot m^{(k-1)} + (1 - \beta) \cdot \left. \frac{\partial L}{\partial \theta} \right|_{\theta=\theta^{(k-1)}}$$
$$\theta^{(k)} \longleftarrow \theta^{(k-1)} - \eta \cdot m^{(k)}$$


Momentum SGD



- ▶ Idea: introduce inertia (like a ball rolling down a hill)
 - Smooths out fast oscillations

$$m^{(k)} \leftarrow \beta \cdot m^{(k-1)} + (1 - \beta) \cdot \left. \frac{\partial L}{\partial \theta} \right|_{\theta=\theta^{(k-1)}}$$
$$\theta^{(k)} \leftarrow \theta^{(k-1)} - \eta \cdot m^{(k)}$$


Momentum SGD



- ▶ Idea: introduce inertia (like a ball rolling down a hill)
 - Smooths out fast oscillations
 - Helps getting out of small local minima

$$m^{(k)} \longleftarrow \beta \cdot m^{(k-1)} + (1 - \beta) \cdot \left. \frac{\partial L}{\partial \theta} \right|_{\theta=\theta^{(k-1)}}$$
$$\theta^{(k)} \longleftarrow \theta^{(k-1)} - \eta \cdot m^{(k)}$$

Momentum SGD



- ▶ Idea: introduce inertia (like a ball rolling down a hill)
 - Smooths out fast oscillations
 - Helps getting out of small local minima
 - Allows for larger range of learning rates*

$$m^{(k)} \leftarrow \beta \cdot m^{(k-1)} + (1 - \beta) \cdot \left. \frac{\partial L}{\partial \theta} \right|_{\theta=\theta^{(k-1)}}$$
$$\theta^{(k)} \leftarrow \theta^{(k-1)} - \eta \cdot m^{(k)}$$

* <https://distill.pub/2017/momentum/>

RMSprop

- ▶ Idea: adjust learning rate separately for different components of the parameter vector
 - Gradients getting smaller \Rightarrow increase the learning rate (scale by inverse running RMS of the gradient)


RMSprop

- ▶ Idea: adjust learning rate separately for different components of the parameter vector
 - Gradients getting smaller \Rightarrow increase the learning rate (scale by inverse running RMS of the gradient)

$$\mathbb{E}[g^2]_{(k)} \leftarrow \beta \cdot \mathbb{E}[g^2]_{(k-1)} + (1 - \beta) \cdot \left(\frac{\partial L}{\partial \theta} \right)^2 \bigg|_{\theta=\theta^{(k-1)}}$$
$$\theta^{(k)} \leftarrow \theta^{(k-1)} - \frac{\eta}{\sqrt{\mathbb{E}[g^2]_{(k)} + \varepsilon}} \cdot \frac{\partial L}{\partial \theta} \bigg|_{\theta=\theta^{(k-1)}}$$

Adam

- ▶ Combine both ideas
(momentum + RMSprop)
- ▶ Typically a good first choice for
an optimizing algorithm



NN generalization



Why deep neural nets generalize well?


- ▶ Number of parameters is often well above the size of the training dataset
- ▶ Would expect heavy overfitting according to “classical ML” theory
- ▶ In practice, test error often decreases with the size of the model

Deep Double Descent

- ▶ In fact, the dependence of the test error from the model size is more complicated
- ▶ Often, the effect of **double descent** is observed
- ▶ Not understood well

- See this review
- Probably, cannot be explained by the implicit regularization from the optimization technique (see, e.g., 2109.14119, 2104.14421)
- Moreover: happens in simpler models, like linear regression (2109.02355)

Img source: <https://openai.com/blog/deep-double-descent/>



Summary

- ▶ Neural networks are essentially **stacked linear models** with scalar nonlinearities in between

Summary

- ▶ Neural networks are essentially **stacked linear models** with scalar nonlinearities in between
- ▶ Earlier layers extract useful features s.t. the problem **becomes solvable** with a linear model (the last layer)

Summary

- ▶ Neural networks are essentially **stacked linear models** with scalar nonlinearities in between
- ▶ Earlier layers extract useful features s.t. the problem **becomes solvable** with a linear model (the last layer)
- ▶ Neural networks can approximate any function arbitrarily well, given they are deep/wide enough

Summary

- ▶ Neural networks are essentially **stacked linear models** with scalar nonlinearities in between
- ▶ Earlier layers extract useful features s.t. the problem **becomes solvable** with a linear model (the last layer)
- ▶ Neural networks can approximate any function arbitrarily well, given they are deep/wide enough
- ▶ Loss functions typically become highly **non-convex** for neural networks
 - this makes the optimization process harder

Summary

- ▶ Neural networks are essentially **stacked linear models** with scalar nonlinearities in between
- ▶ Earlier layers extract useful features s.t. the problem **becomes solvable** with a linear model (the last layer)
- ▶ Neural networks can approximate any function arbitrarily well, given they are deep/wide enough
- ▶ Loss functions typically become highly **non-convex** for neural networks
 - this makes the optimization process harder
- ▶ A variety of **SGD modifications** are available to mitigate this problem

Summary

- ▶ Neural networks are essentially **stacked linear models** with scalar nonlinearities in between
- ▶ Earlier layers extract useful features s.t. the problem **becomes solvable** with a linear model (the last layer)
- ▶ Neural networks can approximate any function arbitrarily well, given they are deep/wide enough
- ▶ Loss functions typically become highly **non-convex** for neural networks
 - this makes the optimization process harder
- ▶ A variety of **SGD modifications** are available to mitigate this problem
- ▶ Food for thought: being the 'universal approximators', can neural nets really solve every possible supervised learning problem?

Thank you!



amaevskij@hse.ru



SiLiKhon

Artem Maevskiy