

Network Regularization

Weight initialization, dropout, batch normalization

Machine Learning and Data Mining, 2022

Artem Maevskiy

National Research University Higher School of Economics

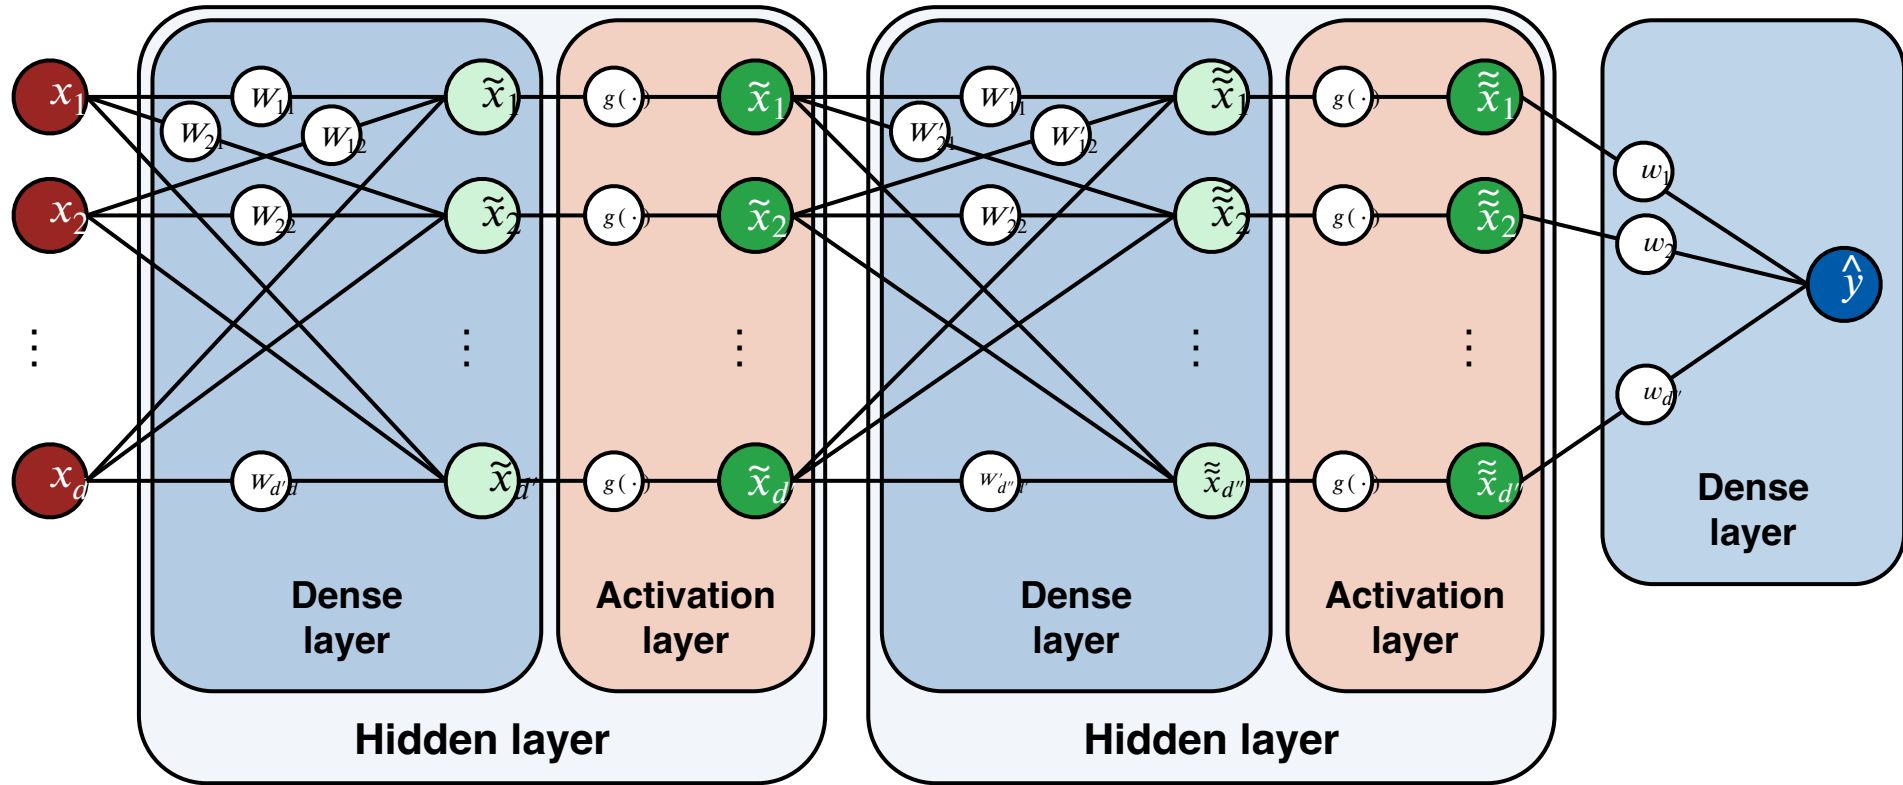


November 17, 2021

Why care about weight initialization?



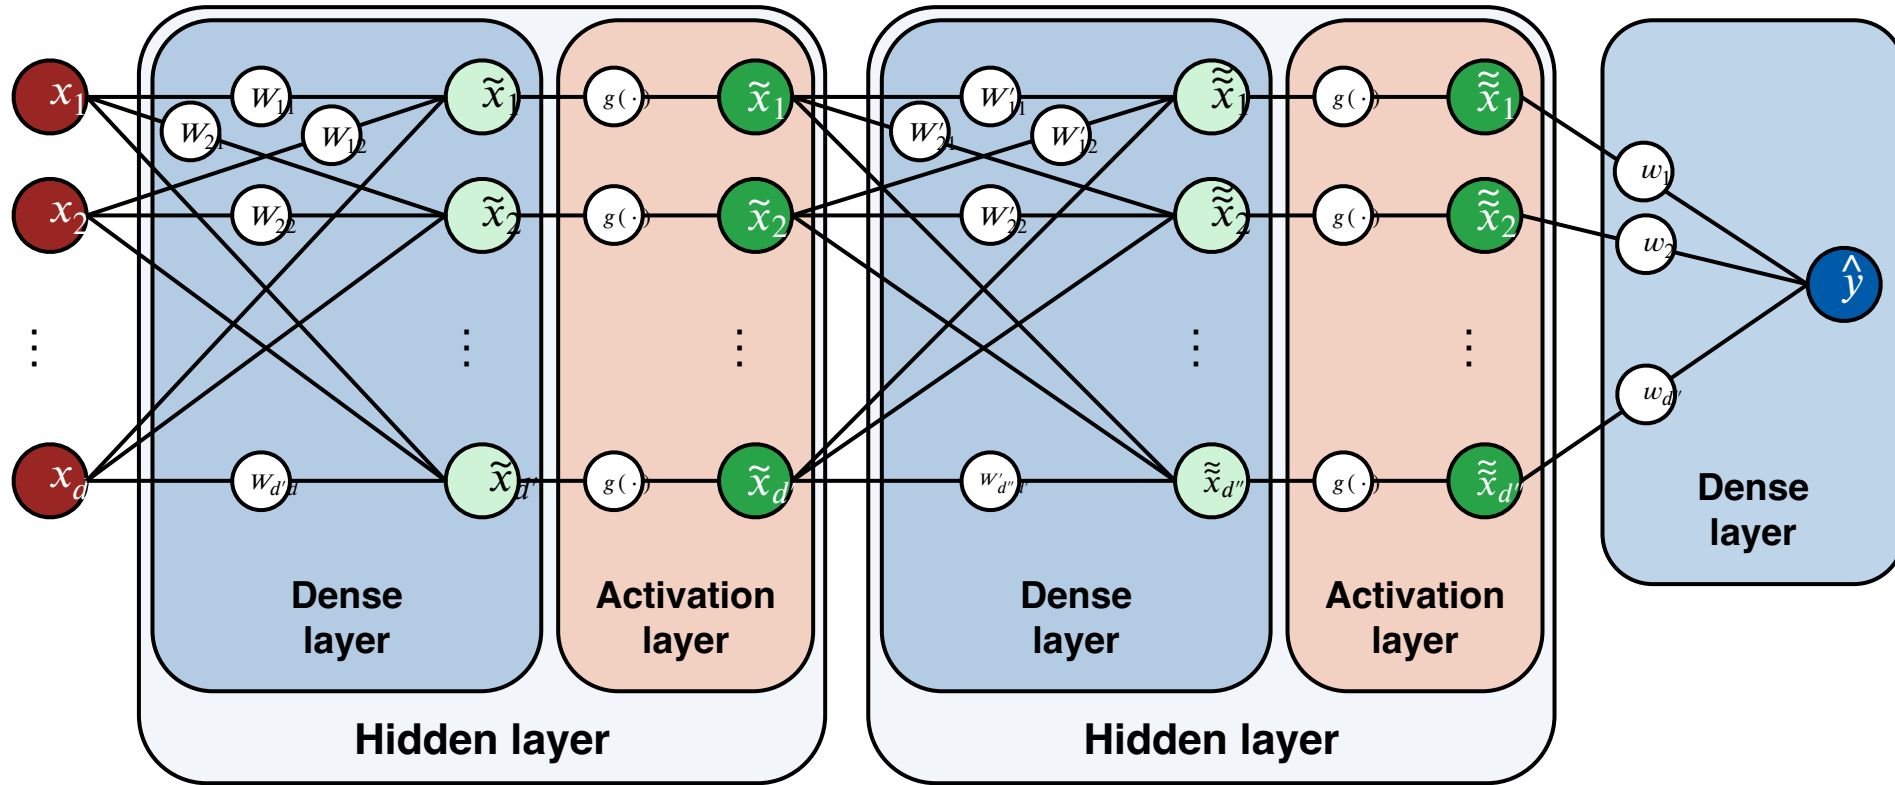
Initialization with a constant (?)



- What happens if we initialize all weights with the same value?

the
is!

Initialization with a constant (?)



- ▶ What happens if we initialize all weights with the same value?
- ▶ Within each layer, the gradients for each of the weights will be the same as well \Rightarrow **updates will be the same** \Rightarrow network degrades!

Initialization with a constant (?)

- ▶ Ok, so constant initialization is a bad idea
- ▶ So, any random initialization should be fine, right?

Some intuition

- ▶ For simplicity, let's omit the activation functions for now
- ▶ Then, the output of a neural network composed of dense layers only is:

$$\hat{y} = W_{out} \cdot \dots \cdot W_{h2} \cdot W_{h1}x$$

Some intuition

- ▶ For simplicity, let's omit the activation functions for now
- ▶ Then, the output of a neural network composed of dense layers only is:

$$\hat{y} = W_{out} \cdot \dots \cdot W_{h2} \cdot W_{h1}x$$

- ▶ Note that gradient wrt to any of the weight matrices W_{hk} is proportional to the **product** of all other matrices

Some intuition

- ▶ For simplicity, let's omit the activation functions for now
- ▶ Then, the output of a neural network composed of dense layers only is:

$$\hat{y} = W_{out} \cdot \dots \cdot W_{h2} \cdot W_{h1}x$$

- ▶ Note that gradient wrt to any of the weight matrices W_{hk} is proportional to the **product** of all other matrices
- ▶ E.g. for 1×1 matrices, if all are of scale $S \in \mathbb{R}$, the gradient g is proportional to:

$$g \sim S^{m-1}$$

where m is the **depth** of the network

Some intuition

- ▶ For simplicity, let's omit the activation functions for now
- ▶ Then, the output of a neural network composed of dense layers only is:

$$\hat{y} = W_{out} \cdot \dots \cdot W_{h2} \cdot W_{h1}x$$


- ▶ Note that gradient wrt to any of the weight matrices W_{hk} is proportional to the **product** of all other matrices
- ▶ E.g. for 1×1 matrices, if all are of scale $S \in \mathbb{R}$, the gradient g is proportional to:

$$g \sim S^{m-1}$$

where m is the **depth** of the network

- ▶ For S too large, the gradients will **explode**; for S too small, they will **vanish**


Some intuition



- More generally:

$$\frac{\partial \text{Loss}}{\partial W_i} = \frac{\partial \text{Loss}}{\partial \text{Output}_i} \cdot \frac{\partial \text{Dense}_i}{\partial W_i} = \frac{\partial \text{Loss}}{\partial \text{Output}_{i+1}} \cdot \frac{\partial \text{Layer}_{i+1}}{\partial \text{Output}_i} \cdot \text{Output}_{i-1}$$

Some intuition




- More generally:

$$\frac{\partial \text{Loss}}{\partial W_i} = \frac{\partial \text{Loss}}{\partial \text{Output}_i} \cdot \frac{\partial \text{Dense}_i}{\partial W_i} = \underbrace{\frac{\partial \text{Loss}}{\partial \text{Output}_{i+1}} \cdot \frac{\partial \text{Layer}_{i+1}}{\partial \text{Output}_i}}_{\text{This will accumulate the product of the gradients for the subsequent layers}} \cdot \text{Output}_{i-1}$$

This will accumulate the product of the gradients for the subsequent layers

Some intuition



- More generally:


$$\frac{\partial \text{Loss}}{\partial W_i} = \frac{\partial \text{Loss}}{\partial \text{Output}_i} \cdot \frac{\partial \text{Dense}_i}{\partial W_i} = \underbrace{\frac{\partial \text{Loss}}{\partial \text{Output}_{i+1}} \cdot \frac{\partial \text{Layer}_{i+1}}{\partial \text{Output}_i}}_{\text{This will accumulate the product of the gradients for the subsequent layers}} \cdot \text{Output}_{i-1}$$

This will accumulate the product of the gradients for the subsequent layers

- Idea: for stable learning we would like to “keep” the scale of the gradients at each step:

$$\text{Var}\left(\frac{\partial \text{Layer}_{i+1}}{\partial \text{Output}_i} \cdot \frac{\partial \text{Layer}_i}{\partial \text{Output}_{i-1}}\right) \approx \text{Var}\left(\frac{\partial \text{Layer}_{i+1}}{\partial \text{Output}_i}\right)$$

Some intuition



- ▶ Similarly, we would also like to not scale the outputs at each step of the forward pass:

$$\text{Var}\left(\text{Layer}_{i+1}\left(\text{Layer}_i\left(\text{Output}_{i-1}\right)\right)\right) \approx \text{Var}\left(\text{Layer}_i\left(\text{Output}_{i-1}\right)\right)$$

Random initialization

$$\begin{aligned} \text{Var}\left(\frac{\partial \mathbf{Layer}_{i+1}}{\partial \mathbf{Output}_i} \cdot \frac{\partial \mathbf{Layer}_i}{\partial \mathbf{Output}_{i-1}}\right) &\approx \text{Var}\left(\frac{\partial \mathbf{Layer}_{i+1}}{\partial \mathbf{Output}_i}\right) \\ \text{Var}\left(\mathbf{Layer}_{i+1}\left(\mathbf{Layer}_i\left(\mathbf{Output}_{i-1}\right)\right)\right) &\approx \text{Var}\left(\mathbf{Layer}_i\left(\mathbf{Output}_{i-1}\right)\right) \end{aligned}$$

Random initialization

$$\text{Var}\left(\frac{\partial \mathbf{Layer}_{i+1}}{\partial \mathbf{Output}_i} \cdot \frac{\partial \mathbf{Layer}_i}{\partial \mathbf{Output}_{i-1}}\right) \approx \text{Var}\left(\frac{\partial \mathbf{Layer}_{i+1}}{\partial \mathbf{Output}_i}\right)$$

$$\text{Var}\left(\mathbf{Layer}_{i+1}\left(\mathbf{Layer}_i\left(\mathbf{Output}_{i-1}\right)\right)\right) \approx \text{Var}\left(\mathbf{Layer}_i\left(\mathbf{Output}_{i-1}\right)\right)$$

- ▶ Generally, these two requirements may contradict each other
- ▶ E.g. for ReLU activation they result in initialization requirements, respectively:

$$\text{Var}\left(W_{ij}\right) = \frac{2}{(\# \text{ outgoing connections})}$$

$$\text{Var}\left(W_{ij}\right) = \frac{2}{(\# \text{ incoming connections})}$$

Random initialization

$$\text{Var}\left(\frac{\partial \text{Layer}_{i+1}}{\partial \text{Output}_i} \cdot \frac{\partial \text{Layer}_i}{\partial \text{Output}_{i-1}}\right) \approx \text{Var}\left(\frac{\partial \text{Layer}_{i+1}}{\partial \text{Output}_i}\right)$$
$$\text{Var}\left(\text{Layer}_{i+1}\left(\text{Layer}_i\left(\text{Output}_{i-1}\right)\right)\right) \approx \text{Var}\left(\text{Layer}_i\left(\text{Output}_{i-1}\right)\right)$$

- ▶ Generally, these two requirements may contradict each other
- ▶ E.g. for ReLU activation they result in initialization requirements, respectively:

$$\text{Var}(W_{ij}) = \frac{2}{(\# \text{ outgoing connections})}$$
$$\text{Var}(W_{ij}) = \frac{2}{(\# \text{ incoming connections})}$$

- ▶ Typically you can just choose one of them, or alternatively average them out:

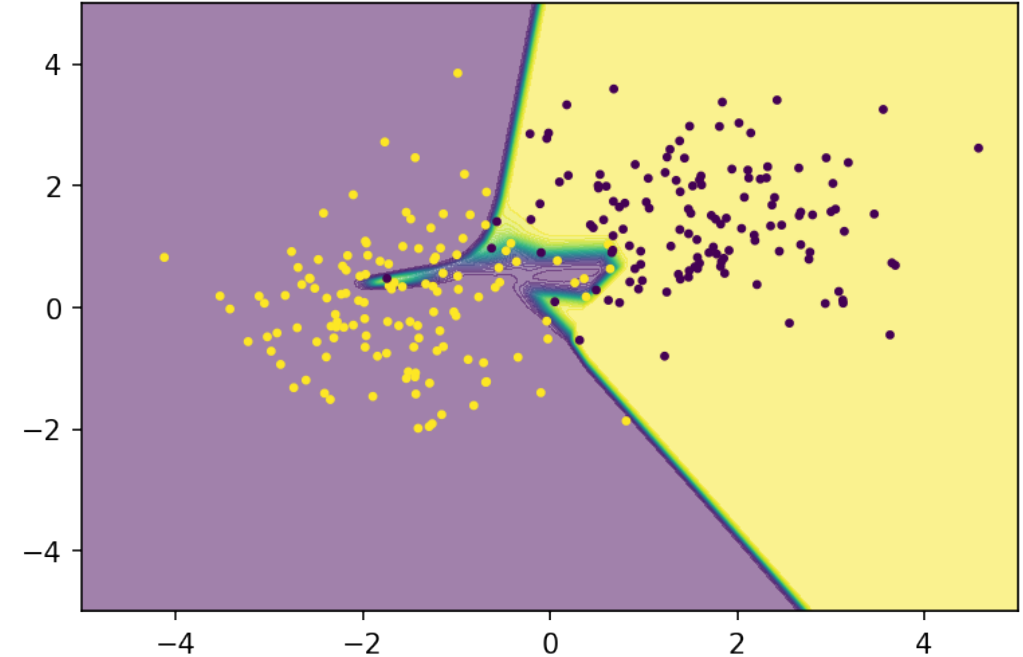

$$\text{Var}(W_{ij}) = \frac{4}{(\# \text{ outgoing connections}) + (\# \text{ incoming connections})}$$

Overfitting with neural networks



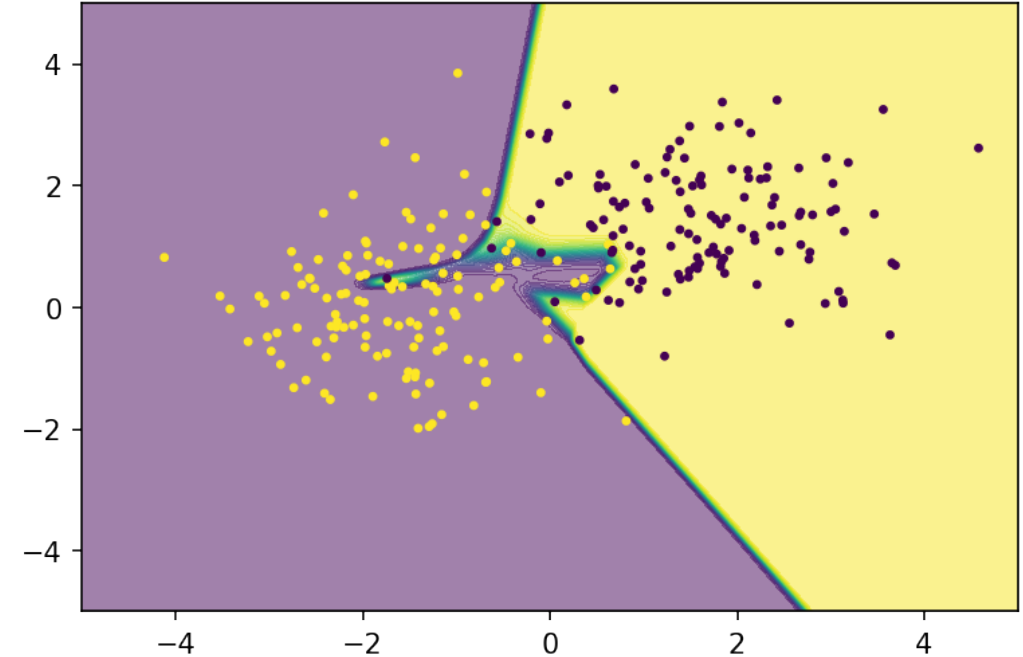

The problem of overfitting

- Being highly complex models, neural networks are prone to overfitting



The problem of overfitting

- ▶ Being highly complex models, neural networks are prone to overfitting

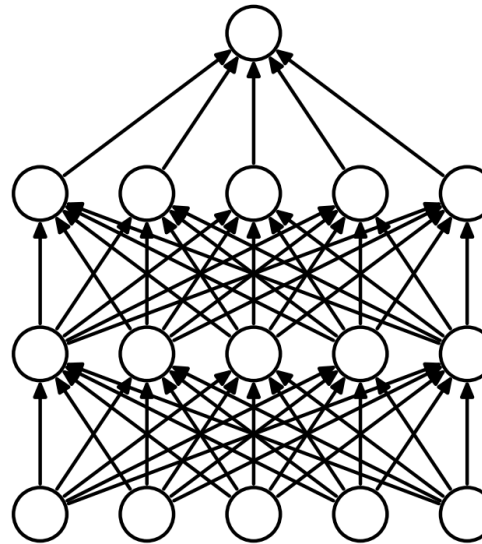


- ▶ Regularization techniques like L1/L2 regularization are also available for neural networks
- ▶ We also discussed **early stopping** (i.e. stop the training before validation error grows)

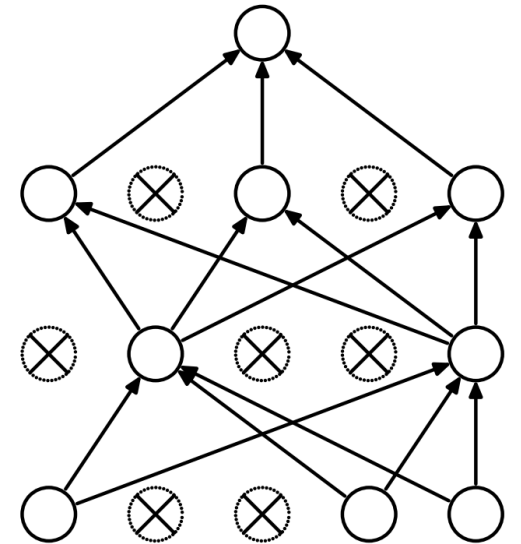
Dropout

- ▶ At train time – sets neuron activations to 0 with a given probability p

Image from:
<http://jmlr.org/papers/v15/srivastava14a.html>



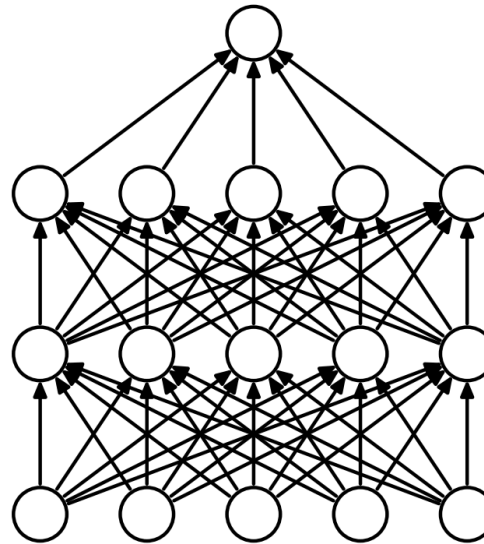
(a) Standard Neural Net



(b) After applying dropout.


Dropout

- ▶ At train time – sets neuron activations to 0 with a given probability p
- ▶ At test time – multiplies the activation by p
 - i.e. sets it to the **expected value**



(a) Standard Neural Net


Image from:
<http://jmlr.org/papers/v15/srivastava14a.html>



(b) After applying dropout.

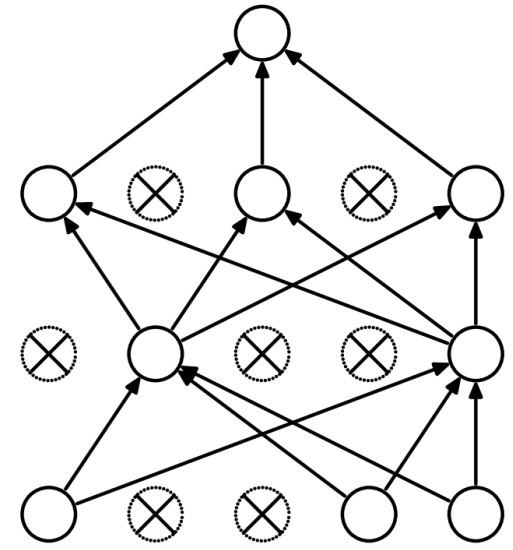
Dropout

- ▶ At train time – sets neuron activations to 0 with a given probability p
- ▶ At test time – multiplies the activation by p
 - i.e. sets it to the **expected value**
- ▶ Makes neuron learn to work with a **randomly chosen sample** of other neurons



(a) Standard Neural Net

Image from:
<http://jmlr.org/papers/v15/srivastava14a.html>

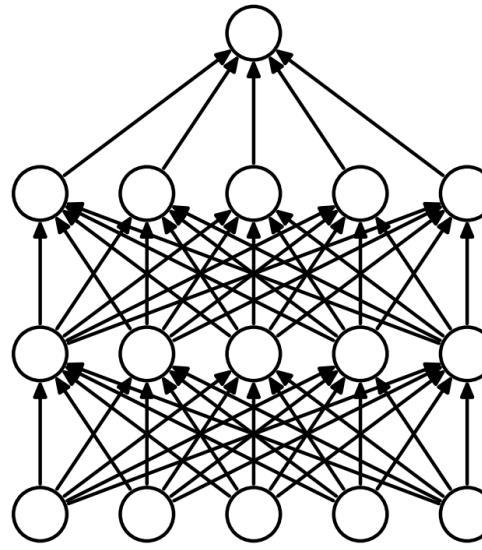


(b) After applying dropout.

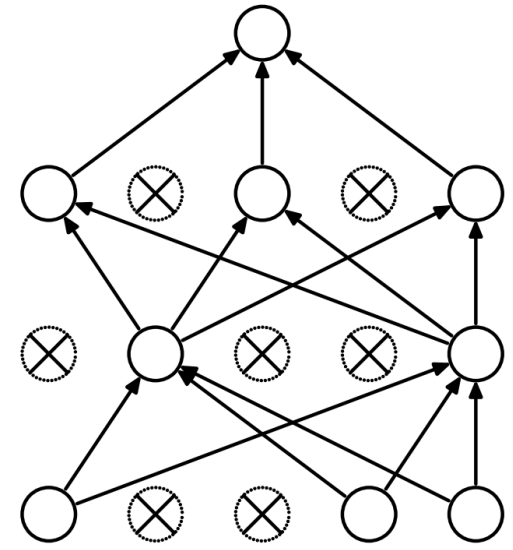
Dropout

- ▶ At train time – sets neuron activations to 0 with a given probability p
- ▶ At test time – multiplies the activation by p
 - i.e. sets it to the **expected value**
- ▶ Makes neuron learn to work with a **randomly chosen sample** of other neurons
- ▶ Drives it towards **creating useful features** rather than relying on other neurons to correct its mistakes

Image from:
<http://jmlr.org/papers/v15/srivastava14a.html>







(a) Standard Neural Net



(b) After applying dropout.

Example from before



In this example, dropout results in a much better (though still not perfect) fit with lower test error

Normalization layers



Batch normalization

- ▶ This technique was originally proposed to mitigate the “internal covariate shift”

internal covariate shift

the updates in one layer
change the input distributions
of the subsequent layers

Batch normalization

- ▶ This technique was originally proposed to mitigate the “internal covariate shift”
- ▶ Works as follows (layer inputs x_i , outputs y_i):
 - calculate sample **mean** and **variance** of the input on a single batch

B

$$\mu_B = \frac{1}{|B|} \sum_{i \in B} x_i$$

$$\sigma_B^2 = \frac{1}{|B|} \sum_{i \in B} (x_i - \mu_B)^2$$

internal covariate shift

the updates in one layer
change the input distributions
of the subsequent layers

Batch normalization

- ▶ This technique was originally proposed to mitigate the “internal covariate shift”
- ▶ Works as follows (layer inputs x_i , outputs y_i):
 - calculate sample **mean** and **variance** of the input on a single batch

B

$$\mu_B = \frac{1}{|B|} \sum_{i \in B} x_i$$

$$\sigma_B^2 = \frac{1}{|B|} \sum_{i \in B} (x_i - \mu_B)^2$$

- **normalize** the input, then **scale and shift** (with the trainable parameters γ , β):

$$y_i = \gamma \cdot \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} + \beta$$

internal covariate shift

the updates in one layer
change the input distributions
of the subsequent layers

Batch normalization

- ▶ Turned out to be **extremely powerful** in many cases
 - Faster and more stable convergence

internal covariate shift

the updates in one layer
change the input distributions
of the subsequent layers

Batch normalization

- ▶ Turned out to be **extremely powerful** in many cases
 - Faster and more stable convergence
- ▶ Later was proved to **not** reduce the internal covariate shift

internal covariate shift

the updates in one layer
change the input distributions
of the subsequent layers

Batch normalization

- ▶ Turned out to be **extremely powerful** in many cases
 - Faster and more stable convergence
- ▶ Later was proved to **not** reduce the internal covariate shift

internal covariate shift

the updates in one layer
change the input distributions
of the subsequent layers

- ▶ Effectively **removes** the ‘shift’ and ‘scale’ degrees of freedom from the previous layer

$$y_i = \gamma \cdot \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} + \beta$$


Batch normalization

- ▶ Which dimension to normalize over? Typically, like this:
 - Batch of 1D vectors [Batch_dim x Features_dim]
 - separately for each component in Features_dim , i.e., over Batch_dim

Batch normalization

- ▶ Which dimension to normalize over? Typically, like this:
 - Batch of 1D vectors [Batch_dim x Features_dim]
 - separately for each component in Features_dim , i.e., over Batch_dim
 - Batch of ND objects [Batch_dim x Spacial_dim1 x ... x Channel_dim]
 - separately for each component in Channel_dim , i.e., over Batch_dim x Spacial_dim1 x ...

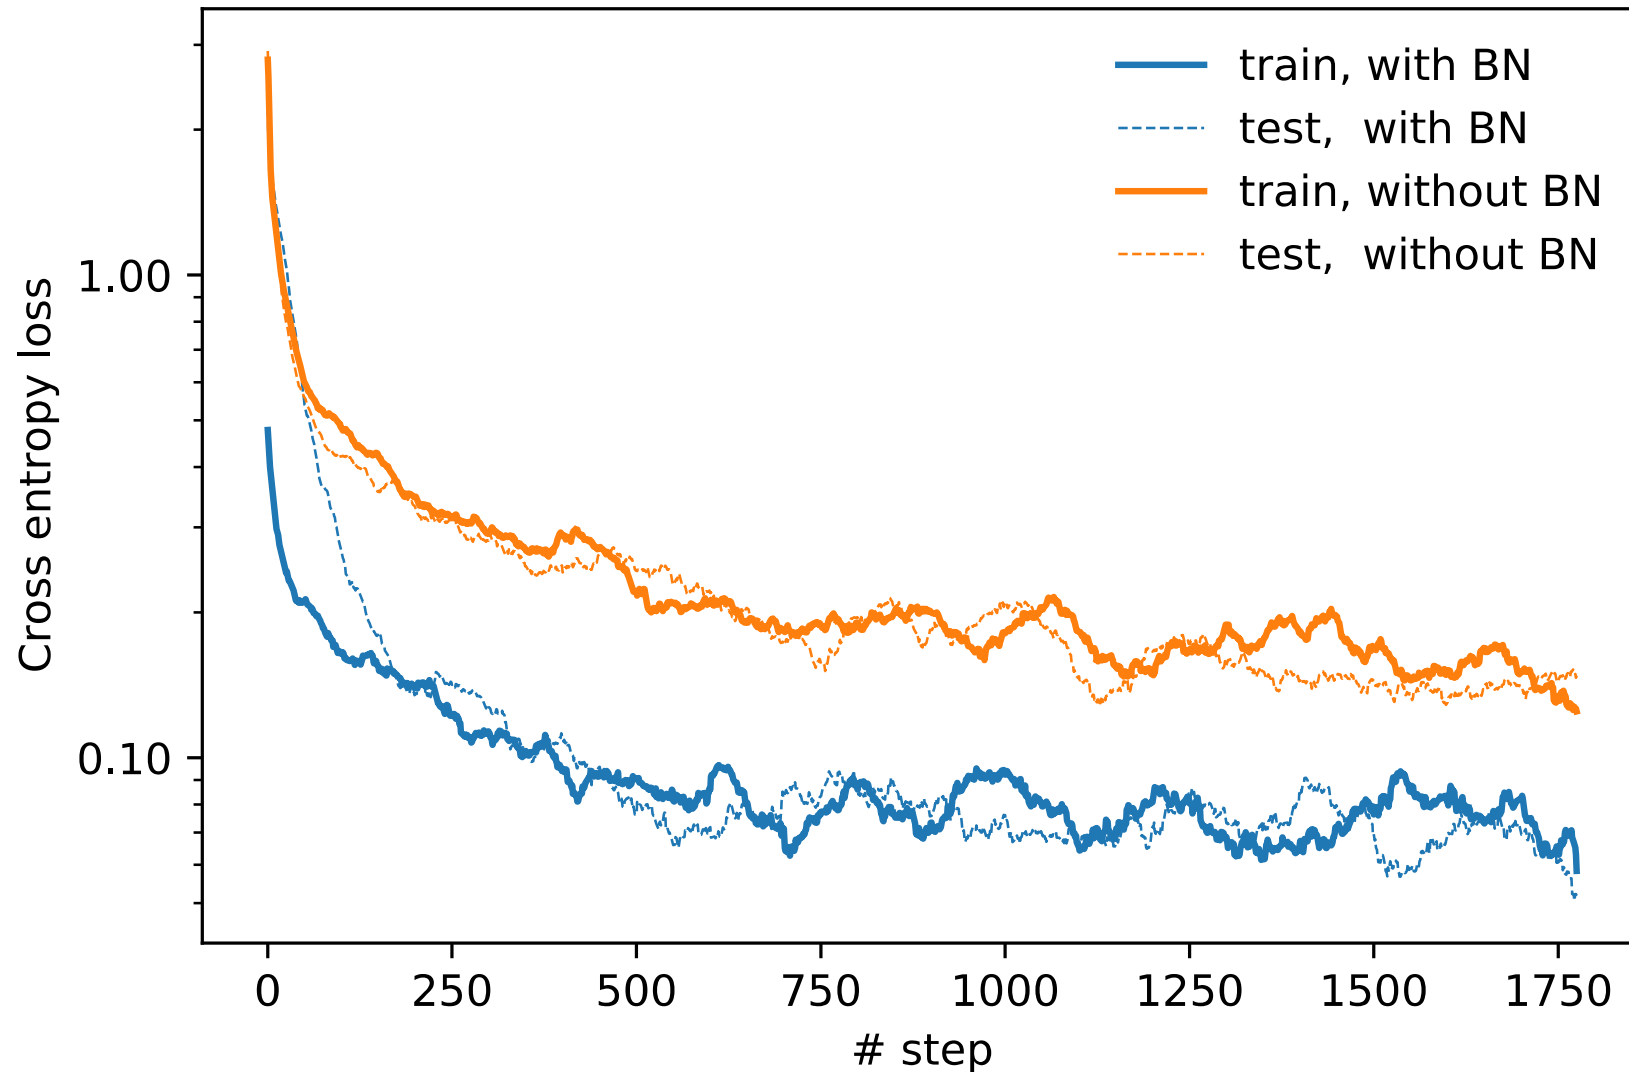
Batch normalization



Batch normalization at inference time

- ▶ Calculating batch statistics at test time may be problematic
 - e.g. when there's a single object to predict
- ▶ Instead: calculate running mean and variance during training, apply at test time

Example: CNN on MNIST



(shown: moving average loss)


Layer Normalization

- ▶ Batch normalization imposes limits on the batch size
 - if too small, the variance of the sample statistics will be too high
- ▶ Problematic to use in recurrent networks

Layer Normalization

- ▶ Batch normalization imposes limits on the batch size
 - if too small, the variance of the sample statistics will be too high
- ▶ Problematic to use in recurrent networks
- ▶ Alternative: **Layer Normalization**
 - the math is same, except statistics is calculated over channels rather than batch elements


Image from <https://arxiv.org/pdf/1803.08494.pdf>



Layer Normalization

- ▶ Batch normalization imposes limits on the batch size
 - if too small, the variance of the sample statistics will be too high
- ▶ Problematic to use in recurrent networks
- ▶ Alternative: **Layer Normalization**
 - the math is same, except statistics is calculated over channels rather than batch elements
 - the effect is quite different though
 - e.g. Layer Normalization “entangles” different neurons within a layer

Image from <https://arxiv.org/pdf/1803.08494.pdf>



Summary

- ▶ If done wrong, weight initialization may cause the gradients to **vanish or explode**

Summary

- ▶ If done wrong, weight initialization may cause the gradients to **vanish or explode**
- ▶ Neural networks can be regularized with L1/L2 penalties or early stopping

Summary

- ▶ If done wrong, weight initialization may cause the gradients to **vanish or explode**
- ▶ Neural networks can be regularized with L1/L2 penalties or early stopping
- ▶ Dropout makes neurons **create useful features** rather than rely on other neurons to correct their mistakes

Summary

- ▶ If done wrong, weight initialization may cause the gradients to **vanish or explode**
- ▶ Neural networks can be regularized with L1/L2 penalties or early stopping
- ▶ Dropout makes neurons **create useful features** rather than rely on other neurons to correct their mistakes
- ▶ Batch normalization is an **extremely powerful** regularization technique, though the reason for that is not entirely clear

Summary

- ▶ If done wrong, weight initialization may cause the gradients to **vanish or explode**
- ▶ Neural networks can be regularized with L1/L2 penalties or early stopping
- ▶ Dropout makes neurons **create useful features** rather than rely on other neurons to correct their mistakes
- ▶ Batch normalization is an **extremely powerful** regularization technique, though the reason for that is not entirely clear
- ▶ Food for thought: how exactly would you implement an early stopping rule?

Thank you!



amaevskij@hse.ru



SiLiKhon

Artem Maevskiy